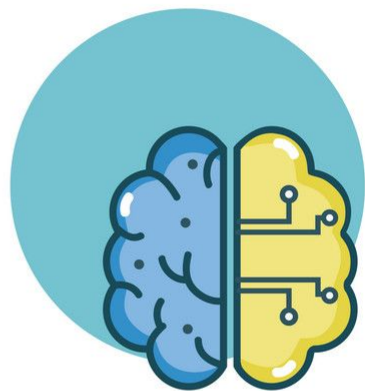


# INTRODUCTION TO MACHINE LEARNING


## K-NEAREST NEIGHBOR



Elisa Ricci



# RECAP: DATA



# RECAP: REPRESENTING EXAMPLES

examples



What is an example?

How is it represented?

# RECAP: FEATURES

examples




features

$f_1, f_2, f_3, \dots, f_n$

$f_1, f_2, f_3, \dots, f_n$

$f_1, f_2, f_3, \dots, f_n$


$f_1, f_2, f_3, \dots, f_n$



How our algorithms  
actually “view” the data

# RECAP: FEATURES

examples



features

red, round, leaf, 3oz, ...

green, round, no leaf, 4oz, ...

yellow, curved, no leaf, 8oz, ...

green, curved, no leaf, 7oz, ...

How our algorithms actually “view” the data

Features are the questions we can ask about the examples

Features in general are represented with vectors

# APPLES VS. BANANAS



Weight	Color	Label
4	Red	Apple
5	Yellow	Apple
6	Yellow	Banana
3	Red	Apple
7	Yellow	Banana
8	Yellow	Banana
6	Yellow	Apple


Can we visualize this data?

# APPLES VS. BANANAS

We can turn features into numerical values

peso in grammi

Weight	Color	Label
4	0	Apple
5	1	Apple
6	1	Banana
3	0	Apple
7	1	Banana
8	1	Banana
6	1	Apple



Examples are points in an  $n$ -dimensional space where  $n$  is the number of features

# EXAMPLES IN A FEATURE SPACE

We can imagine to have 3 categories of interest



feature<sub>2</sub>


feature<sub>1</sub>

ho solo 2 features e quindi  
uno spazio 2D è sufficiente


- label 1
- label 2
- label 3




# TRAINING SET



TEST EXAMPLE: WHICH CLASS?



# WHAT ABOUT THIS EXAMPLE?



# K-NEAREST NEIGHBOR (K-NN)


To classify an example  $d$ :

- Find  $k$  nearest neighbors of  $d$
- Choose as the label the **majority label** within the  $k$  nearest neighbors

Robust to noisy data by considering  $k$ -nearest neighbors

# EUCLIDEAN DISTANCE


In two dimensions, how do we compute the distance?



$$D(a, b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2}$$

# EUCLIDEAN DISTANCE

In n-dimensions, how do we compute the distance?



$$D(a, b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_n - b_n)^2}$$

# EUCLIDEAN DISTANCE

- Euclidean Distance makes sense when different features are comparable (e.g. each is variable measured in the same units).
  - For instance, if the measurements are different, say length and weight, it is not clear.
- **Standardization:** [data pre-processing](#) When features are not comparable we can standardize them by dividing by their corresponding standard deviation. This makes them all equally important.

# STANDARDIZATION & SCALING

- **Standardization or Z-score normalization**

- Rescale the data so that the mean is 0 and the standard deviation from the mean (standard scores) is 1

$$x_{norm} = \frac{x - \mu}{\sigma}$$

$\mu$  is mean,  $\sigma$  is a standard deviation from the mean  
(standard score)

- **Min-Max scaling**

- Scale the data to a fixed range – between 0 and 1

$$x_{morm} = \frac{x - x_{min}}{x_{max} - x_{min}}$$



# DISTANCE AND SIMILARITY

- Measuring distance/similarity is a domain-specific problem and there are many, many different variations
- Similarity
  - Numerical measure of how alike two data objects are.
  - Is higher when objects are more alike.
  - Often falls in the range  $[0,1]$
- Distance
  - Numerical measure of how different are two data objects
  - Lower when objects are more alike
  - Minimum dissimilarity is often 0
  - Upper limit varies

# MINKOWSKI DISTANCE

- Minkowski Distance is a generalization of Euclidean Distance


$$D(a, b) = \sum_{k=1}^p |a_k - b_k|^r$$

*a, b = vettori*


- $r$  is a parameter,  $p$  is the number of dimensions
- $r = 1$ . City block (Manhattan, L1 norm) distance.
- $r = 2$ . Euclidean distance
- $r \rightarrow \infty$ . “supremum” ( $L_\infty$  norm) distance.
  - This is the maximum difference between any component of the vectors

# DISTANCES


Euclidean



Manhattan



$r = \infty$   
Minkowski




# COSINE SIMILARITY

Cosine Similarity

Given two document vectors:

$$\cos(d_1, d_2) = (d_1 \cdot d_2) / (||d_1|| ||d_2||),$$

Where  $\cdot$  indicates vector dot product and  $||d||$  is the length of vector  $d$ .



Example:

$$d_1 = 3 \ 2 \ 0 \ 5 \ 0 \ 0 \ 0 \ 2 \ 0 \ 0$$


$$d_2 = 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 2$$

$$d_1 \cdot d_2 = 3*1 + 2*0 + 0*0 + 5*0 + 0*0 + 0*0 + 0*0 + 2*1 + 0*0 + 0*2 = 5$$

$$||d_1|| = (3*3 + 2*2 + 0*0 + 5*5 + 0*0 + 0*0 + 0*0 + 2*2 + 0*0 + 0*0)^{0.5} = (42)^{0.5} = 6.481$$

$$||d_2|| = (1*1 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 1*1 + 0*0 + 2*2)^{0.5} = (6)^{0.5} = 2.245$$

$$\cos(d_1, d_2) = .3150$$




DECISION BOUNDARIES

HYPER-PARAMETERS

# DECISION BOUNDARIES


The **decision boundaries** are places in the features space where the classification of a point/example changes



# DECISION BOUNDARIES


Voronoi

k-NN gives **locally** defined decision boundaries between classes




# VORONOI DIAGRAM

- Describes the areas that are nearest to any given point, given a set of data.
- Each line segment is equidistant between two points



**(a)** Initial points distribution




**(b)** Voronoi diagram by initial points




# DECISION BOUNDARIES

- k-NN algorithm does not explicitly compute decision boundaries.
- The decision boundaries form a subset of the Voronoi diagram for the training data.
- The more examples that are stored, the more complex the decision boundaries can become




# CHOOSING K

What is the label with  $k = 3$ ?



## CHOOSING K


What is the label with  $k = 20$ ?






# CHOOSING K

k deve essere sempre minore di n(numero di punti)

What is the label with  $k = 20$ ? **Red!** (class with max number of datapoints, weird)



# THE IMPACT OF K



What is the role of  $k$ ?

How does it relate to overfitting and underfitting?

# RECAP: UNDERFITTING OVERFITTING

- The factors determining how well a machine learning algorithm will perform are its ability to:
  - Make the **training error small**.
  - Make the **gap between training and test error small**.
- These two factors correspond to the cases of **underfitting** and **overfitting**.
- **Underfitting** occurs when the model is not able to obtain a sufficiently low error value on the training set.
- **Overfitting** occurs when the gap between the training error and test error is too large.

# HOW TO PICK K

- Common heuristics:
  - often 3, 5, 7
  - choose an odd number to avoid ties
- Use validation set
- Use cross-validation
- Rule of thumb is  $k < \sqrt{n}$ , where  $n$  is the number of training examples

# K-NN VARIATIONS

Instead of  $k$  nearest neighbors, count majority from all examples within a fixed distance

## Weighted $k$ -NN: [weighted points](#)

- Right now, all examples are treated equally
- Weight the “vote” of the examples, so that closer examples have more vote/weight
- Often use some sort of exponential decay






# LAZY LEARNER VS EAGER LEARNER

no training  
processing only in the  
inference phase

- k-NN is belongs to the class of **lazy learning** algorithms
- **Lazy learning**: Simply stores training data (or perform limited processing) and operates when it is given a test example
- **Eager learning** (e.g. Decision Trees, SVMs): Given a training set, constructs a classification model before receiving new test data to classify
- **Lazy learners**: less time in training but more time in predicting

# CURSE OF DIMENSIONALITY

- **Curse of dimensionality:** In high dimensions almost all points are far away from each other
- The size of the data space grows exponentially with the number of dimensions.
- This means that the size of the data set must also grow exponentially in order to keep the same density.



# CURSE OF DIMENSIONALITY

- The success of k-NN is very dependent on having a **dense data set**.
- Every machine learning algorithm needs a dense data set for accurate prediction
- What makes k-nearest neighbors special? tutte le features sono importanti  
perchè calcolo la distanza in tutte le dimensioni
- k-NN requires a point to be close in **every single dimension**.
- Some algorithms can create models based on single dimensions, and only need points to be close together along that axis.
- k-NN doesn't work that way. It needs all points to be close along every axis in the data space.



# COMPUTATIONAL COST

- Linear algorithm (no pre-processing):
  - Compute distance for all  $N$  datapoints
  - Complexity of distance computation:  $O(kN)$
  - No additional space needed
- Tree-based **data structures**: pre-processing
  - Often used in applications: k-d trees ( k-dimensional trees).

# K-D TREE

l'idea è quella di muoversi un sottogruppo preciso prima di calcolare la "vicinanza" computo questa struttura dati prima per velocizzare i calcoli durante la inference phase

- Building a K-D tree from training data:
  - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
  - pick random dimension, find median, split data, repeat
- Find NNs for new point  $(7,4)$ 
  - find region containing  $(7,4)$
  - compare to all points in region



# AN ASIDE: PARAMETRIC VS NON-PARAMETRIC MODELS

- Parametric models we have a finite number of parameters (curve fitting)  
ex: parametri del polinomio
  - Linear regression, logistic regression, and linear Support Vector Machines
- Nonparametric models: the number of parameters is (potentially) infinite. The complexity of the model grows with the number of training data. (k-nn)
  - K-nearest neighbor, decision trees, or RBF kernel SVMs are considered as non-parametric learning algorithms since the number of parameters grows with the size of the training set.

# K-NEAREST NEIGHBOR: SUMMARY

- Non parametric model
- When to consider
  - Instance map to points in  $R^N$
  - Few (e.g. less than 20) features per instance
  - Lots of training data
- Advantages
  - Training is very fast (no training)
  - Learn complex target functions
- Disadvantages
  - Slow at query time
  - Easily fooled by irrelevant attributes

# K-NEAREST NEIGHBOR: ADVANTAGES


- Easy to program
- No optimization or training required
- Classification accuracy can be very good, it can outperform more complex models

resisted the test of time



# EXAMPLE

Decent performance when large training set



- Yann LeCunn – MNIST Digit Recognition
  - Handwritten digits
  - 28x28 pixel images:  $d = 784$
  - 60,000 training samples
  - 10,000 test samples
- Nearest neighbour is competitive

	Test Error Rate (%)
Linear classifier (1-layer NN)	12.0
K-nearest-neighbors, Euclidean	5.0
K-nearest-neighbors, Euclidean, deskewed	2.4
K-NN, Tangent Distance, 16x16	1.1
K-NN, shape context matching	0.67
1000 RBF + linear classifier	3.6
SVM deg 4 polynomial	1.1
2-layer NN, 300 hidden units	4.7
2-layer NN, 300 HU, [deskewing]	1.6
LeNet-5, [distortions]	0.8
Boosted LeNet-4, [distortions]	0.7

specialized distance

# K-NEAREST NEIGHBOR: ISSUES

- Choose distance measure
  - Most common: Euclidean
- Choosing  $k$
- Curse of Dimensionality
- Memory based technique.
  - Must make a pass through the data for each classification. This can be prohibitive for large data sets.

# QUESTIONS?

