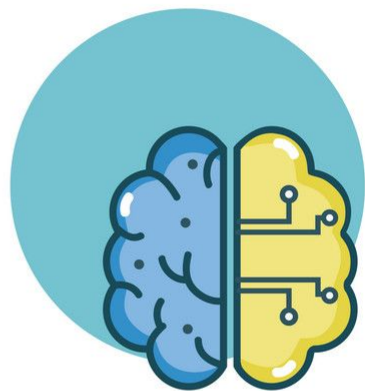



# INTRODUCTION TO MACHINE LEARNING

## GRADIENT DESCENT



Elisa Ricci






# MODELS AND ALGORITHMS

# MACHINE LEARNING IDEA


- ML allows computers to acquire knowledge.
- Knowledge is acquired through **algorithms** by learning and inferring from **data**.
- Knowledge is represented by a **model**.
- The model is used on future data.



# LINEAR MODELS

A **linear model** is a model assumes that the data are linearly separable

Assume a specific hypothesis space, i.e. linear functions



# LINEAR MODELS

A linear model in  $n$ -dimensional space (i.e.  $n$  features) is defined by  $n+1$  weights.

In two dimensions, we have a line:

$$0 = w_1 f_1 + w_2 f_2 + b$$

In three dimensions, a plane:

$$0 = w_1 f_1 + w_2 f_2 + w_3 f_3 + b$$

In  $n$ -dimensions, a **hyperplane**

$$0 = b + \sum_{i=1}^n w_i f_i$$

# PERCEPTRON LEARNING ALGORITHM

**repeat** until convergence (or for some # of iterations):

**for** each training example ( $f_1, f_2, \dots, f_n$ , label):

$$prediction = b + \sum_{i=1}^n w_i f_i$$


**if**  $prediction * label \leq 0$ : // they don't agree

**for** each  $w_i$ :

$$w_i = w_i + f_i * label$$

$$b = b + label$$

# WHICH LINE WILL THE PERCEPTRON FIND?



Only guaranteed to find **some** line that separates the data!

# LINEAR MODELS

Perceptron algorithm is one example of a linear classifier

Many, many other algorithms learn a line (i.e. a setting of a linear combination of weights)

Goals:

- Explore a number of linear training algorithms
- Understand *why these algorithms work*



# PERCEPTRON LEARNING ALGORITHM

**repeat** until convergence (or for some # of iterations):

**for** each training example ( $f_1, f_2, \dots, f_n$ , label):

$$prediction = b + \sum_{i=1}^n w_i f_i$$

**if**  $prediction * label \leq 0$ : // they don't agree

**for** each  $w_i$ :

$$w_i = w_i + f_i * label$$

$$b = b + label$$

# A CLOSER LOOK AT WHY WE GOT IT WRONG

$$w_1 \quad w_2$$
$$1 * f_1 + 0 * f_2 =$$

$$1 * -1 + 0 * 1 = -1$$

$$(-1, 1) \quad +$$

← This value should be positive!value

↑  
contributed in the  
wrong direction

←  
could have contributed  
(positive feature) but it did  
not since the weight is 0

# A CLOSER LOOK AT WHY WE GOT IT WRONG

$$w_1 \quad w_2$$
$$1 * f_1 + 0 * f_2 =$$

$$1 * -1 + 0 * 1 = -1$$

$$(-1, 1) \quad +$$

← This value should be positive!value

↑  
decrease

e.g. from 1 to 0

← increase

from 0 to 1

# MODEL-BASED MACHINE LEARNING

## 1. Pick a model

- e.g. a hyperplane, a decision tree,...
- A model is defined by a collection of parameters

What are the parameters for DT? Perceptron?

# MODEL-BASED MACHINE LEARNING

## 1. Pick a model

- e.g. a hyperplane, a decision tree,...
- A model is defined by a collection of parameters

DT: the structure of the tree, which features each node splits on, the predictions at the leaves

Perceptron: the weights and the b value

# MODEL-BASED MACHINE LEARNING

1. Pick a model
  - e.g. a hyperplane, a decision tree,...
  - A model is defined by a collection of parameters
2. Pick a criterion to optimize (aka objective function)

What criteria do decision tree learning and perceptron learning optimize?

# MODEL-BASED MACHINE LEARNING

1. Pick a **model**
  - e.g. a hyperplane, a decision tree,...
  - A model is defined by a collection of parameters
2. Pick a criterion to optimize (aka **objective function**)
  - e.g. training error
3. Develop a **learning algorithm**
  - the algorithm should try and minimize the criteria, sometimes in a heuristic way (i.e. non-optimally), sometimes exactly

# LINEAR MODELS IN GENERAL

1. Pick a model

$$0 = \underbrace{b}_{\text{bias}} + \sum_{j=1}^m \underbrace{w_j}_{\text{weights}} f_j$$

These are the parameters we want to learn

2. Pick a criterion to optimize (aka objective function)



# SOME NOTATION: INDICATOR FUNCTION

Convenient notation for turning True and False answers into numbers/counts:

$$1[x] = \begin{cases} 1 & \text{if } x = \text{True} \\ 0 & \text{if } x = \text{False} \end{cases}$$

# SOME NOTATION: DOT-PRODUCT

We use a **vector notation**

We represent an example  $f_1, f_2, \dots, f_m$  as a single vector,  $\mathbf{x}$

- $j$  subscript will indicate feature indexing, i.e.,  $x_j$
- $i$  subscript will indicate examples indexing over a dataset, i.e.,  $x_i$  or sometimes  $x_{ij}$

Similarly, we can represent the weight  $w_1, w_2, \dots, w_m$  as a single vector,  $\mathbf{w}$

The dot-product between two vectors  $\mathbf{a}$  and  $\mathbf{b}$  is defined as:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{j=1}^m a_j b_j$$

# LINEAR MODELS

1. Pick a model

$$0 = b + \sum_{j=1}^n w_j f_j$$

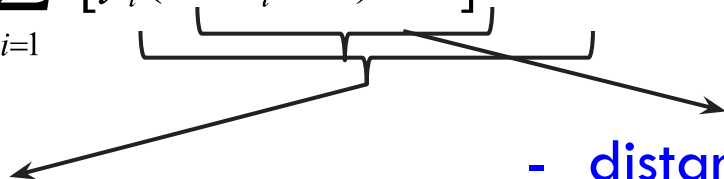
These are the parameters we want to learn

2. Pick a criterion to optimize (aka objective function)

$$\sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$

What does this equation say?

# 0/1 LOSS FUNCTION

$$\sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$


whether or not the  
prediction and label agree,  
true if *they don't*

- distance from hyperplane
- sign is prediction

total number of mistakes,  
aka 0/1 loss

# MODEL-BASED MACHINE LEARNING

1. Pick a model

$$0 = b + \sum_{j=1}^m w_j f_j$$


2. Pick a criteria to optimize (aka objective function)

$$\sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$

3. Develop a learning algorithm

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$

Find  $w$  and  $b$  that  
minimize the 0/1 loss  
(i.e. training error)



# LOSS FUNCTIONS

# MINIMIZING 0/1 LOSS

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$

Find  $w$  and  $b$  that  
minimize the 0/1 loss


How do we do this?

How do we *minimize* a function?

Why is it hard for this function?

# MINIMIZING 0/1 IN ONE DIMENSION

$$\sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$




Each time we change  $w$  such that the example is right/wrong the loss will increase/decrease



# MINIMIZING 0/1 IN ONE DIMENSION

$$\sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$



Each new feature we add (i.e. weights) adds another dimension to this space!

# MINIMIZING 0/1 LOSS

Find  $w$  and  $b$  that  
minimize the 0/1 loss


$$\operatorname{argmin}_{w,b} \sum_{i=1}^n 1[y_i(w \cdot x_i + b) \leq 0]$$

This turns out to be hard (in fact, NP-HARD 😞)

Challenge:


- Small changes in any  $w$  can have large changes in the loss (the change isn't continuous)
- There can be many, many local minima
- At any given point, we don't have much information to direct us towards any minima

# MORE MANAGEABLE LOSS FUNCTIONS



What property/properties do we want from our loss function?



# MORE MANAGEABLE LOSS FUNCTIONS



- Ideally, continuous (i.e. differentiable) so we get an indication of direction of minimization
- Only one minima

# CONVEX FUNCTIONS

Convex functions look something like:



One definition: The line segment between any two points on the function is *above* the function

# SURROGATE LOSS FUNCTIONS

For many applications, we really would like to minimize the 0/1 loss

A **surrogate loss function** is a loss function that provides an upper bound on the actual loss function (in this case, 0/1)

We'd like to identify convex surrogate loss functions to make them easier to minimize

**Key to a loss function:** how it scores the difference between the actual label  $y$  and the predicted label  $y'$

# SURROGATE LOSS FUNCTIONS

0/1 loss:

$$l(y, y') = 1[y y' \leq 0]$$

Ideas?

Some function that is a proxy for error, but is continuous and convex

# SURROGATE LOSS FUNCTIONS

0/1 loss:

$$l(y, y') = 1[y y' \leq 0]$$

Hinge:

$$l(y, y') = \max(0, 1 - y y')$$

Exponential:

$$l(y, y') = \exp(-y y')$$

Squared loss:

$$l(y, y') = (y - y')^2$$

Why do these work? What do they penalize?




# SURROGATE LOSS FUNCTIONS

0/1 loss:  $l(y, y') = 1[y y' \leq 0]$

Hinge:  $l(y, y') = \max(0, 1 - y y')$

Squared loss:  $l(y, y') = (y - y')^2$

Exponential:  $l(y, y') = \exp(-y y')$



# MODEL-BASED MACHINE LEARNING

1. pick a model

$$0 = b + \sum_{j=1}^m w_j f_j$$

2. pick a criteria to optimize (aka objective function)


$$\sum_{i=1}^n \exp(-y_i(w \cdot x_i + b))$$

use a convex surrogate  
loss function

3. develop a learning algorithm


$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b))$$

Find  $w$  and  $b$  that  
minimize the  
surrogate loss



# GRADIENT DESCENT


# FINDING THE MINIMUM



How do we find the minimum for a function?

# ONE APPROACH: GRADIENT DESCENT


Partial derivatives give us the slope (i.e. direction to move) in that dimension



# ONE APPROACH: GRADIENT DESCENT

## Approach:


- pick a starting point ( $w$ )
- repeat:
  - pick a dimension
  - move a small amount in that dimension towards decreasing loss (using the derivative)



# ONE APPROACH: GRADIENT DESCENT

## Approach:

- pick a starting point ( $w$ )
- repeat:
  - pick a dimension
  - move a small amount in that dimension towards decreasing loss (using the derivative)



# GRADIENT DESCENT

- Pick a starting point ( $w$ )
- Repeat until loss doesn't decrease in any dimension:
  - pick a dimension
  - move a small amount in that dimension towards decreasing loss (using the derivative)

$$w_j = w_j - \frac{d}{dw_i} \text{loss}(w)$$

Why negative?





# GRADIENT DESCENT

$$w_j = w_j - \eta \frac{d}{dw_j} loss(w)$$

Learning rate

How much we want to move in the error direction, often this will change over time

# SOME MATH

$$\begin{aligned}\frac{d}{dw_j} loss &= \frac{d}{dw_j} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) \\ &= \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) \frac{d}{dw_j} (-y_i(w \cdot x_i + b))\end{aligned}$$

# SOME MATH

$$\begin{aligned} -\frac{d}{dw_i} y_i (w \cdot x_i + b) &= -\frac{d}{dw_i} y_i (\sum_{j=1}^m w_j x_{ij} + b) \\ &= -\frac{d}{dw_i} y_i (w_1 x_{i1} + w_2 x_{i2} + \dots + w_m x_{im} + b) \\ &= -\frac{d}{dw_i} y_i w_1 x_{i1} + y_i w_2 x_{i2} + \dots + y_i w_m x_{im} + y_i b \\ &= -y_i x_{ij} \end{aligned}$$

# SOME MATH

$$\begin{aligned}\frac{d}{dw_j} loss &= \frac{d}{dw_j} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) \\ &= \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) \frac{d}{dw_j} - y_i(w \cdot x_i + b) \\ &= \sum_{i=1}^n -y_i x_{ij} \exp(-y_i(w \cdot x_i + b))\end{aligned}$$

# GRADIENT DESCENT

- For our choice of the loss we have:

$$w_j = w_j - \eta \frac{d}{dw_j} \text{loss}(w)$$

$$w_j = w_j + \eta \sum_{i=1}^n y_i x_{ij} \exp(-y_i(w \cdot x_i + b))$$

What is this doing?

# EXPONENTIAL UPDATE RULE

$$w_j = w_j + \eta \sum_{i=1}^n y_i x_{ij} \exp(-y_i (w \cdot x_i + b))$$

---

for each example  $x_i$ :

$$w_j = w_j + \eta y_i x_{ij} \exp(-y_i (w \cdot x_i + b))$$

Does this look familiar?

# PERCEPTRON LEARNING ALGORITHM!

**repeat** until convergence (or for some # of iterations):

**for** each training example ( $f_1, f_2, \dots, f_n$ , label):

$$\text{prediction} = b + \sum_{i=1}^n w_i f_i$$

**if**  $\text{prediction} * \text{label} \leq 0$ : // they don't agree

**for** each  $w_i$ :

$$w_i = w_i + f_i * \text{label}$$

$$b = b + \text{label}$$

---

$$w_j = w_j + \eta y_i x_{ij} \exp(-y_i(w \cdot x_i + b))$$

In practice  $w_j = w_j + x_{ij} y_i c$       where  $c = \eta \exp(-y_i(w \cdot x_i + b))$

# PERCEPTRON LEARNING ALGORITHM!

**repeat** until convergence (or for some # of iterations):

**for** each training example ( $f_1, f_2, \dots, f_n$ , label):

$$\text{prediction} = b + \sum_{i=1}^n w_i f_i$$

**if**  $\text{prediction} * \text{label} \leq 0$ : // they don't agree

**for** each  $w_i$ :

**Note:** for gradient descent, we always update

$$w_i = w_i + f_i * \text{label}$$

$$b = b + \text{label}$$

$$w_j = w_j + \eta y_i x_{ij} \exp(-y_i(w \cdot x_i + b))$$

In practice  $w_j = w_j + x_{ij} y_i c$  where  $c = \eta \exp(-y_i(w \cdot x_i + b))$



# THE CONSTANT

$$c = \eta \exp(-y_i(w \cdot x_i + b))$$

learning rate

label

prediction

When is this large/small?

# THE CONSTANT

$$c = \eta \exp(-y_i(w \cdot x_i + b))$$

label

prediction


- If they are the same sign, as the predicted gets larger there update gets smaller
- If they are different, the more different they are, the bigger the update

# GRADIENT

- The gradient is the vector of partial derivatives wrt to all the coordinates of the weights:

$$\nabla_{\mathbf{w}} L = \left[ \frac{\partial L}{\partial w_1} \quad \frac{\partial L}{\partial w_2} \quad \cdots \quad \frac{\partial L}{\partial w_N} \right]$$

- Each partial derivative measures how fast the loss changes in one direction.
- When the gradient is zero, i.e. all the partials derivatives are zero, the loss is not changing in any direction.
- Note: the arrows (gradients) point out from a minimum toward a maximum.



# GRADIENT DESCENT

---


**Algorithm 21** GRADIENTDESCENT( $\mathcal{F}, K, \eta_1, \dots$ )

---

```
1:  $\mathbf{z}^{(0)} \leftarrow \langle o, o, \dots, o \rangle$  // initialize variable we are optimizing
2: for  $k = 1 \dots K$  do
3:    $\mathbf{g}^{(k)} \leftarrow \nabla_{\mathbf{z}} \mathcal{F}|_{\mathbf{z}^{(k-1)}}$  // compute gradient at current location
4:    $\mathbf{z}^{(k)} \leftarrow \mathbf{z}^{(k-1)} - \eta^{(k)} \mathbf{g}^{(k)}$  // take a step down the gradient
5: end for
6: return  $\mathbf{z}^{(K)}$ 
```


---

# GRADIENT DESCENT




# GRADIENT DESCENT


- In problems where the optimization problem is non-convex, it probably has **local minima**.
- An example is neural network. For long time people did not use NN, because they prefer methods that are guarantees to find the optimal solution.



# GRADIENT DESCENT




# GRADIENT DESCENT





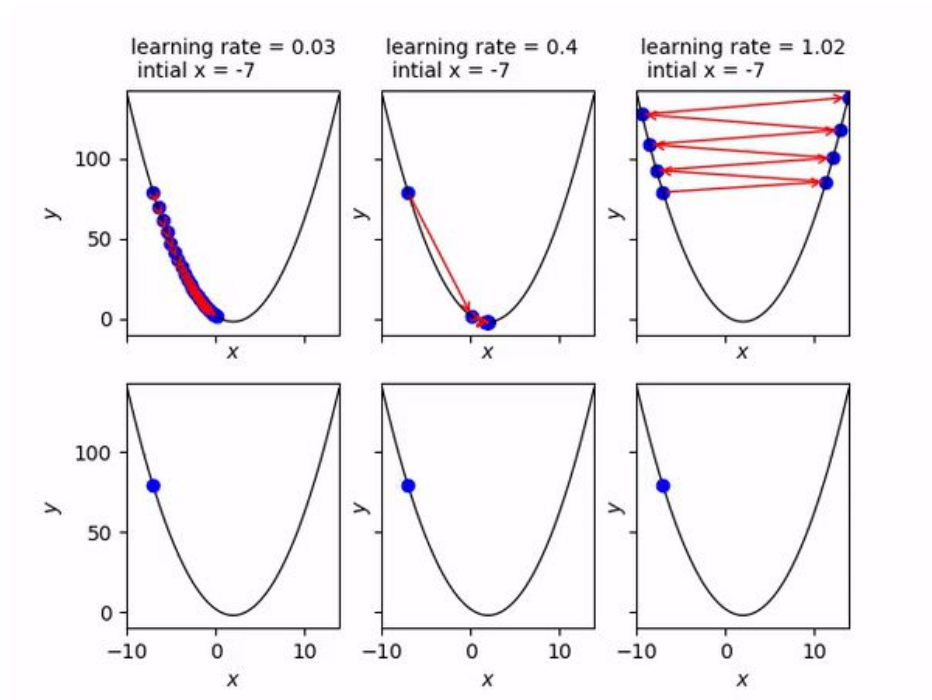
# GRADIENT DESCENT

- **Saddle point:** Some directions curve upwards, and others curve downwards.
- At a saddle point, the gradient is 0 even if we are not at a minimum.
- If we are exactly on the saddle point, then we are stuck.
- If we are slightly to the side, then we can get unstuck.
- **Saddle points very common in high dimensions!**



# LEARNING RATE

- Very important hyper-parameter



# SUMMARY

- Model-based machine learning:
  - define a model, objective function (i.e. loss function), minimization algorithm
- Gradient descent minimization algorithm
  - so far we consider the case where the loss function is convex
  - make small updates towards lower losses
- Perceptron learning algorithm and gradient descent/exponential loss function (modulo a learning rate)
- Gradient descent in general

# QUESTIONS?



Some slides are taken from David Kauchak