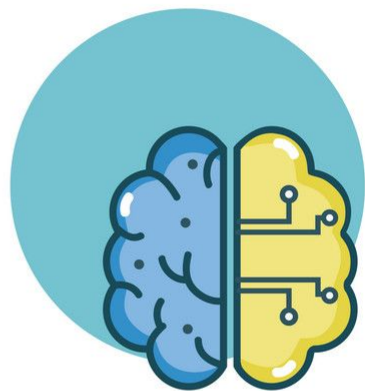



# INTRODUCTION TO MACHINE LEARNING

## REGULARIZATION



Elisa Ricci





AVOID OVERFITTING


# ONE CONCERN

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b))$$

We are calculating this on the **training set**

We still need to be careful about **overfitting!**

The  $\min_{w,b} \text{Loss}$  on the training set is generally NOT the min for the test set



How did we deal with this?

# REGULARIZATION

- A **regularizer** is an additional criterion to the loss function to make sure that we do not overfit
- It is called a regularizer since it tries to keep the parameters more normal/regular
- It is a bias on the model that forces the learning to prefer certain types of weights over others


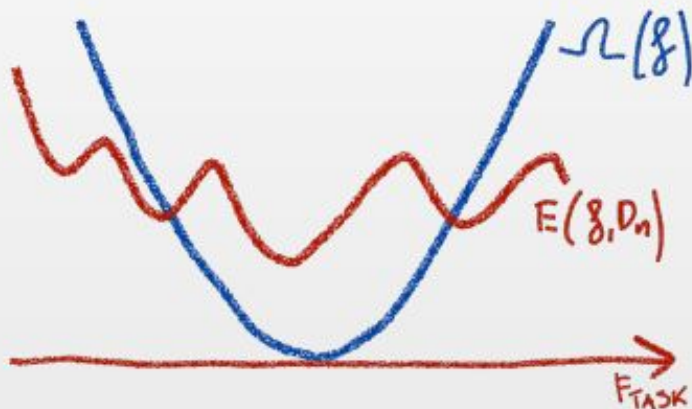
$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \operatorname{loss}(y y') + \lambda \operatorname{regularizer}(w, b)$$


# SO FAR...

Modification of the training error function with a term  $\Omega(f)$  that typically penalizes complex solutions

$$E_{\text{reg}}(f; \mathcal{D}_n) = E(f; \mathcal{D}_n) + \lambda_n \Omega(f)$$

TRADE-OFF  
PARAMETER





# REGULARIZERS

# REGULARIZERS

$$0 = b + \sum_{j=1}^n w_j f_j$$

- Generally, we do not want huge weights: if weights are large, a small change in a feature can result in a large change in the prediction
- Might also prefer weights of 0 for features that are not useful

# COMMON REGULARIZERS

sum of the weights

$$r(w, b) = \sum |w_j|$$

sum of the squared weights

$$r(w, b) = \sqrt{\sum |w_j|^2}$$

Squared weights penalizes large values more  
Sum of weights will penalize small values more



# P-NORM

sum of the weights (1-norm)

$$r(w, b) = \sum |w_j|$$

sum of the squared weights  
(2-norm)

$$r(w, b) = \sqrt{\sum |w_j|^2}$$

---

p-norm = family of norms


p-norm

$$r(w, b) = \sqrt[p]{\sum |w_j|^p} = \|w\|^p$$

Smaller values of p ( $p < 2$ ) encourage sparser vectors  
Larger values of p discourage large weights more


# L1/L2-NORMS VISUALIZED

invece di trovare una soluzione al centro come prima ora  
la trovo nell'intersezione tra la loss function e il  
regularizer



il quadrato è il regularizer


(a)  $\ell_1$ -ball meets quadratic function.  
 $\ell_1$ -ball has corners. It's very likely that  
the meet-point is at one of the corners.



il cerchio è il regularizer

(b)  $\ell_2$ -ball meets quadratic function.  
 $\ell_2$ -ball has no corner. It is very unlikely  
that the meet-point is on any of axes."

# P-NORMS VISUALIZED



all  $p$ -norms penalize larger weights

$p < 2$  tends to create sparse  
(i.e. lots of 0 weights)

$p > 2$  tends to like similar  
weights

# MODEL-BASED MACHINE LEARNING

1. pick a model

$$0 = b + \sum_{j=1}^n w_j f_j$$

2. pick a criteria to optimize (aka objective function)

$$\sum_{i=1}^n \text{loss}(yy') + \lambda \text{regularizer}(w)$$

non c'è una regola per scegliere  $\lambda$   
se non trial and error  
è uno dei tanti hyperparameter

3. develop a learning algorithm

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \text{loss}(yy') + \lambda \text{regularizer}(w)$$

Find  $w$  and  $b$   
that minimize

# MINIMIZING WITH A REGULARIZER

We know how to solve convex minimization problems using gradient descent:

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \operatorname{loss}(yy')$$

If we can ensure that the loss + regularizer is convex then we could still use gradient descent:

$$\operatorname{argmin}_{w,b} \underbrace{\sum_{i=1}^n \operatorname{loss}(yy') + \lambda \operatorname{regularizer}(w)}_{\text{convex as long as both loss and regularizer are convex}}$$

convex as long as both loss and regularizer are convex

# MODEL-BASED MACHINE LEARNING

1. Pick a model

$$0 = b + \sum_{j=1}^n w_j f_j$$

2. Pick a criteria to optimize (aka objective function)

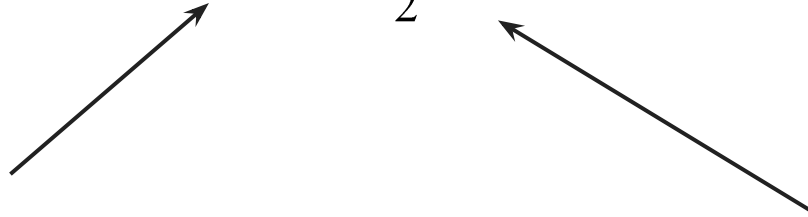
$$\sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \frac{\lambda}{2} \|w\|^2$$

3. Develop a learning algorithm

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \frac{\lambda}{2} \|w\|^2$$

Find  $w$  and  $b$   
that minimize

# OUR OPTIMIZATION CRITERION

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \frac{\lambda}{2} \|w\|^2$$


Loss function: penalizes examples where the prediction is different than the label

Regularizer: penalizes large weights

Key: this function is convex allowing us to use gradient descent

# GRADIENT DESCENT

- pick a starting point ( $w$ )
- repeat until loss doesn't decrease in any dimension:
  - pick a dimension
  - move a small amount in that dimension towards decreasing loss (using the derivative)

$$w_j = w_j - \eta \frac{d}{dw_j} (\text{loss}(w) + \text{regularizer}(w, b))$$

---

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i (w \cdot x_i + b)) + \frac{\lambda}{2} \|w\|^2$$



# SOME MORE MATHS

$$\frac{d}{dw_j} \text{objective} = \frac{d}{dw_j} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \frac{\lambda}{2} \|w\|^2$$

⋮

(some math happens)

$$= - \sum_{i=1}^n y_i x_{ij} \exp(-y_i(w \cdot x_i + b)) + \lambda w_j$$

# GRADIENT DESCENT

- pick a starting point ( $w$ )
- repeat until loss doesn't decrease in any dimension:
  - pick a dimension
  - move a small amount in that dimension towards decreasing loss (using the derivative)

$$w_j = w_j - \eta \frac{d}{dw_j} (\text{loss}(w) + \text{regularizer}(w, b))$$

---

$$w_j = w_j + \eta \sum_{i=1}^n y_i x_{ij} \exp(-y_i (w \cdot x_i + b)) - \eta \lambda w_j$$

# THE UPDATE

$$w_j = w_j + \eta y_i x_{ij} \exp(-y_i (w \cdot x_i + b)) - \eta \lambda w_j$$

learning rate

direction to  
update

constant: how far from wrong

regularization

If  $w_i$  is positive, reduces  $w_i$   
If  $w_i$  is negative, increases  $w_i$

} moves  $w_i$  towards 0  $i \rightarrow j$

# L1 REGULARIZATION

$$\operatorname{argmin}_{w,b} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \|w\|$$

---

$$\frac{d}{dw_j} \sum_{i=1}^n \exp(-y_i(w \cdot x_i + b)) + \lambda \|w\|$$

$$= -\sum_{i=1}^n y_i x_{ij} \exp(-y_i(w \cdot x_i + b)) + \lambda \operatorname{sign}(w_j)$$

# THE UPDATE

$$w_j = w_j + \eta y_i x_{ij} \exp(-y_i (w \cdot x_i + b)) - \eta \lambda \text{sign}(w_j)$$

learning rate

direction to  
update

constant: how far from wrong

regularization

If  $w_i$  is positive, reduces by a constant

If  $w_i$  is negative, increases by a constant

} moves  $w_i$  towards 0  
***regardless of magnitude***

# REGULARIZATION WITH P-NORMS

**L1:**

$$w_j = w_j + \eta(loss\_correction - \lambda sign(w_j))$$

**L2:**

$$w_j = w_j + \eta(loss\_correction - \lambda w_j)$$

**Lp:**

$$w_j = w_j + \eta(loss\_correction - \lambda c w_j^{p-1})$$

# REGULARIZERS SUMMARIZED

- L1 is popular because it tends to result in sparse solutions (i.e. lots of zero weights). However, it is not differentiable, so it only works for gradient descent solvers
- L2 is also popular because for some loss functions, it can be solved directly (no gradient descent required, though often iterative solvers still)
- $L_p$  is less popular since they don't tend to shrink the weights enough

# THE OTHER LOSS FUNCTIONS

Without regularization, the generic update is:

$$w_j = w_j + \eta y_i x_{ij} c$$

where

$$c = \exp(-y_i(w \cdot x_i + b))$$

exponential

$$c = 1[y y' < 1]$$

hinge loss

---

$$w_j = w_j + \eta (y_i - (w \cdot x_i + b)) x_{ij}$$

squared error



# QUESTIONS?



Some slides are taken from David Kauchak