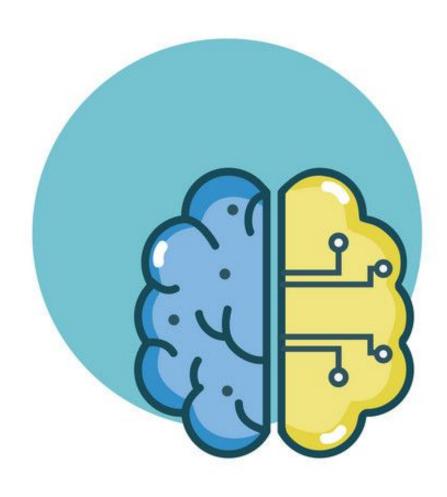
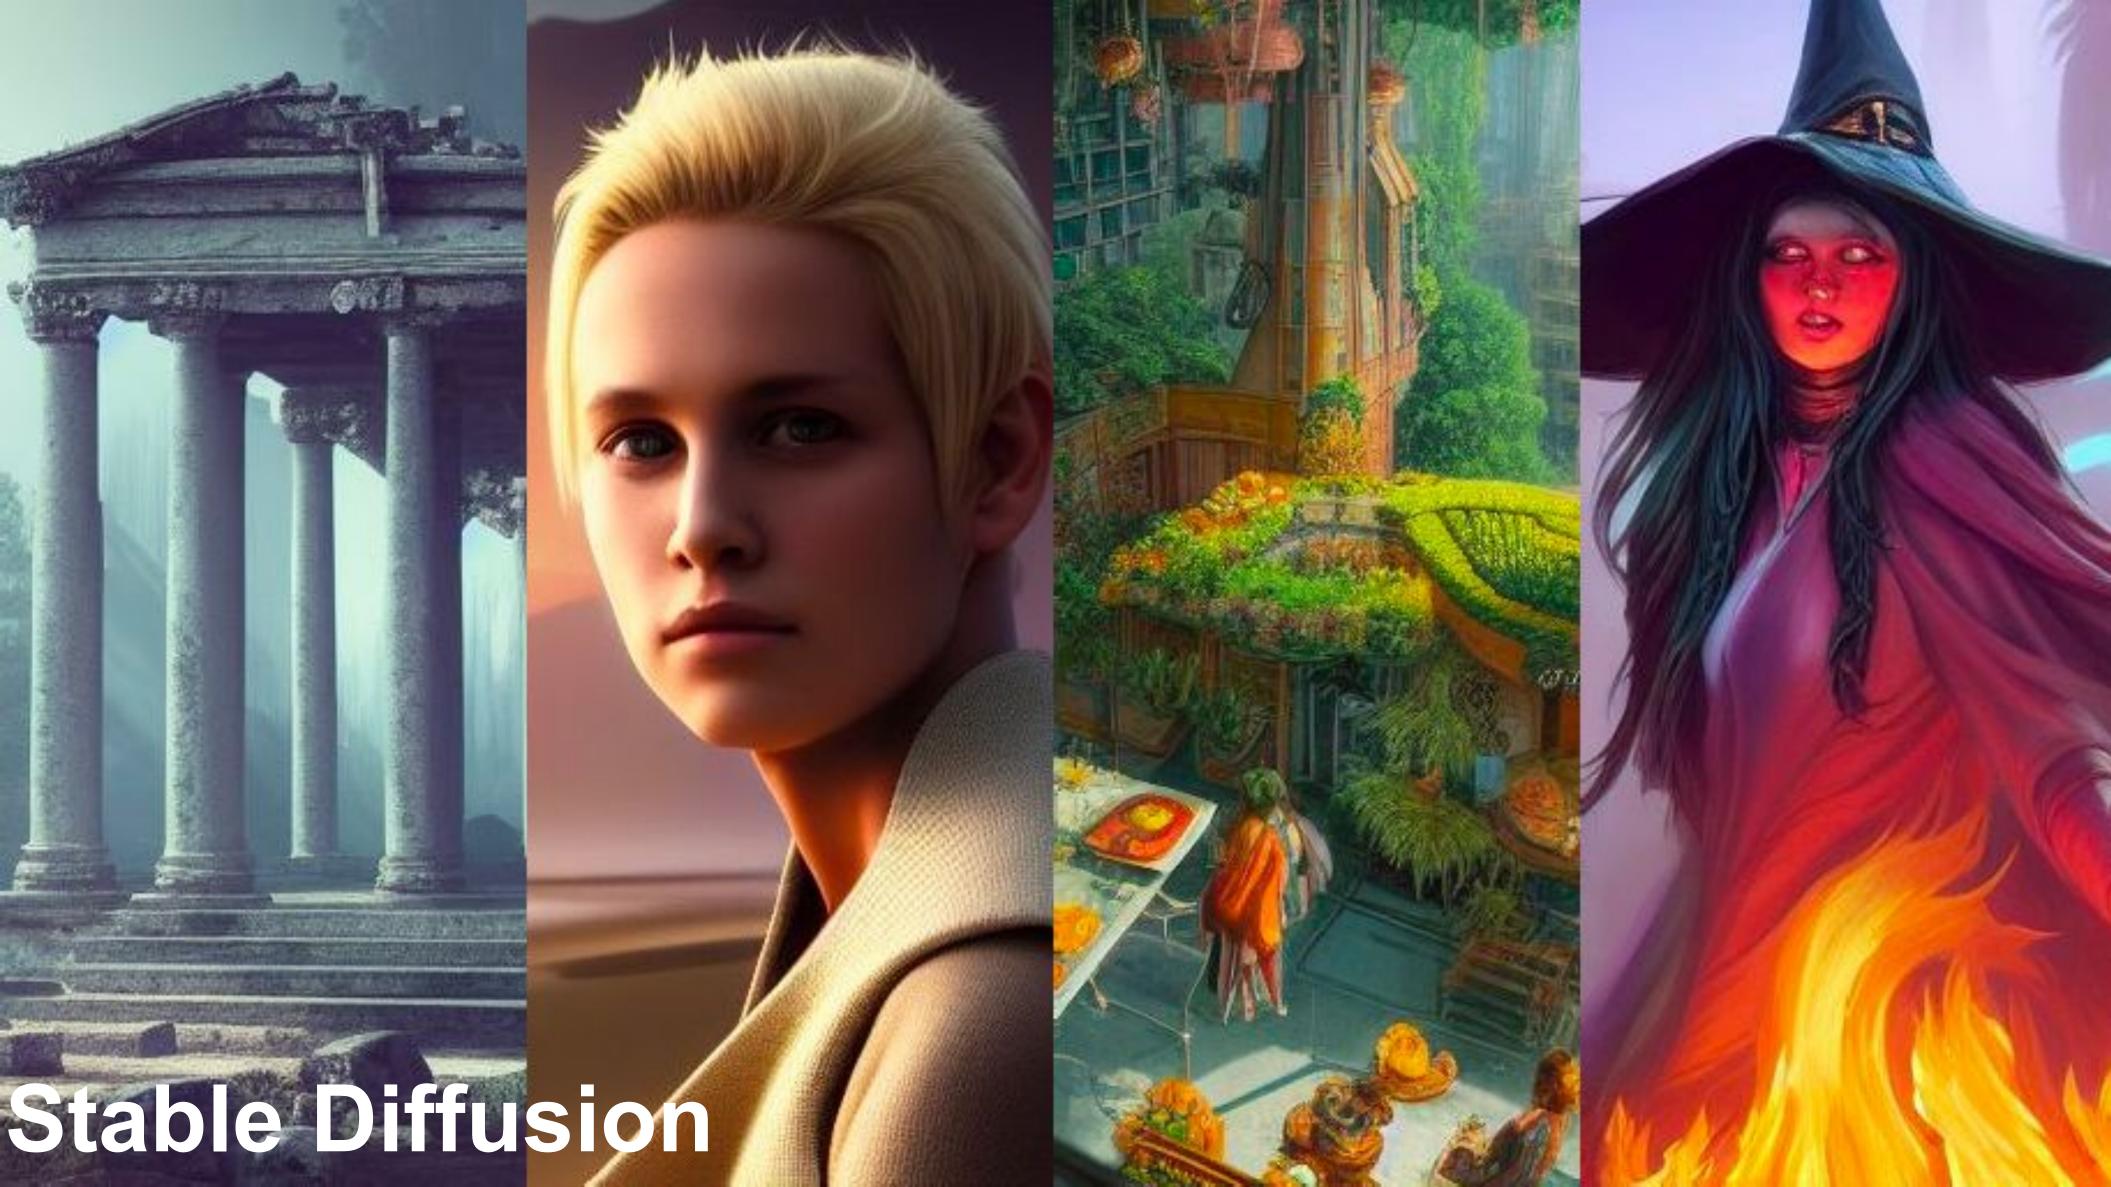
# INTRODUCTION TO MACHINE LEARNING

#### DIFFUSION MODELS



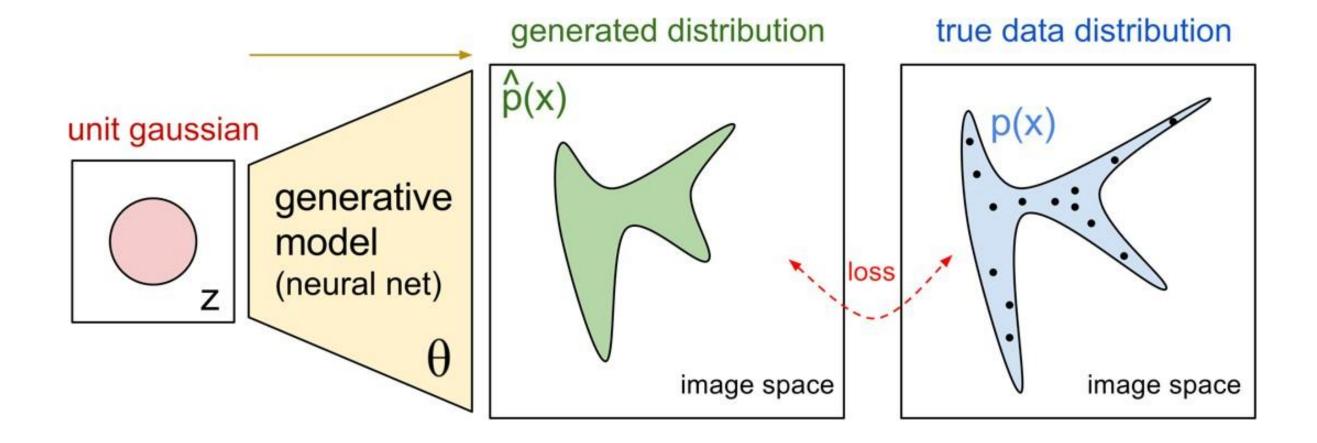
Elisa Ricci





#### GENERATIVE MODELS ricordarsi che ques non solo immagini

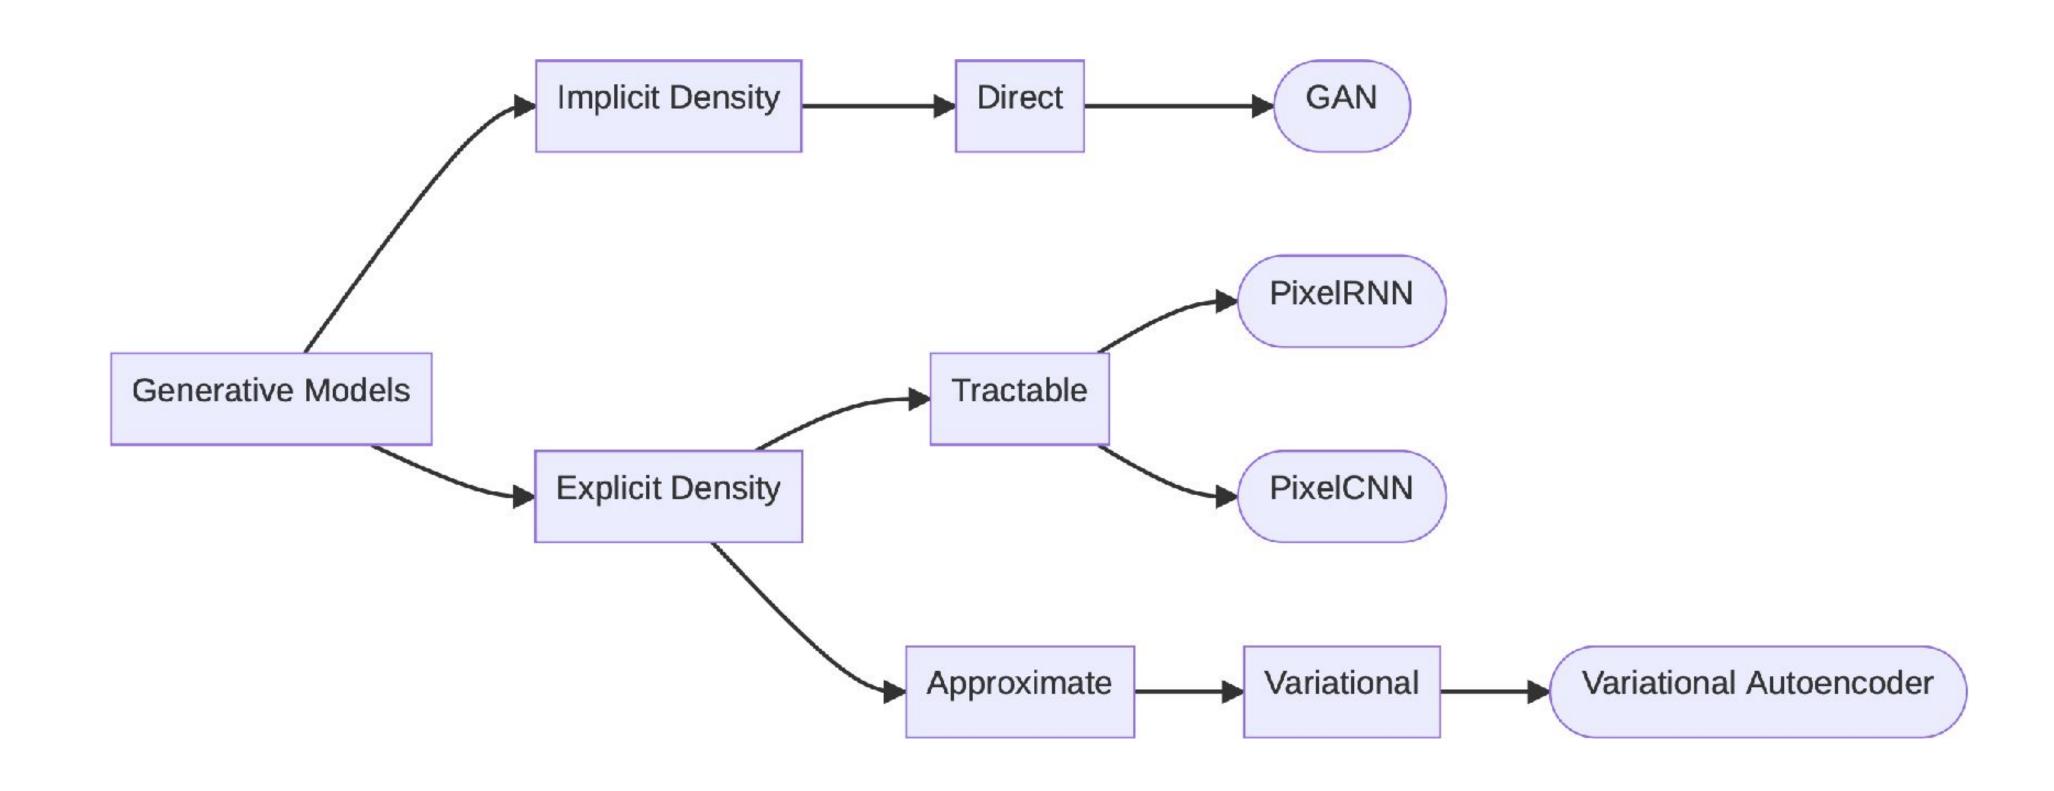
ricordarsi che questi modelli vanno bene per ogni tipo di structured data non solo immagini



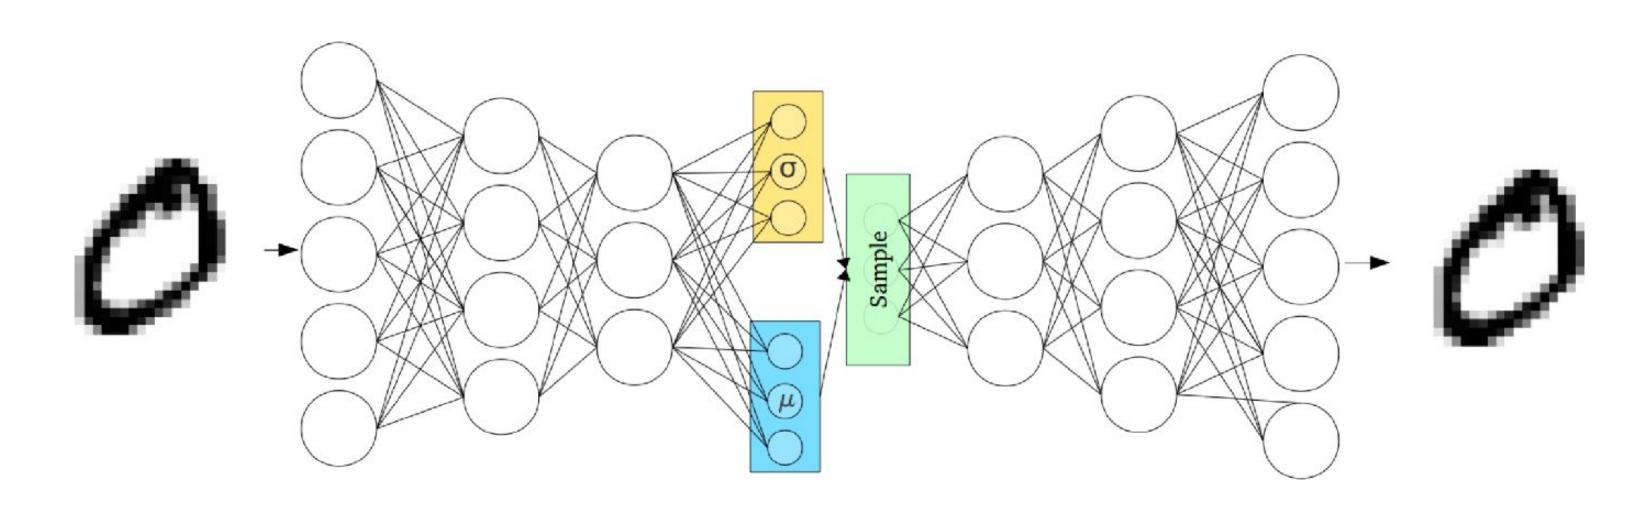
Train a generator G from latent space to data space and approximate the real data distribution Training is difficult:

- Hyperparameters choice
- Quantify similarity between sets
- Choice of latent space

#### GENERATIVE MODELS

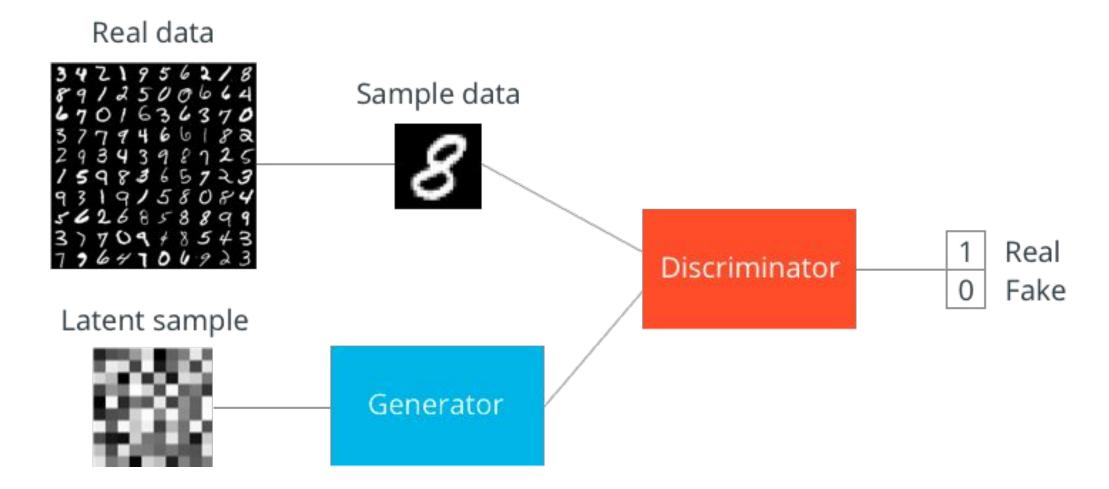


#### VARIATIONAL AUTOENCODERS Recap



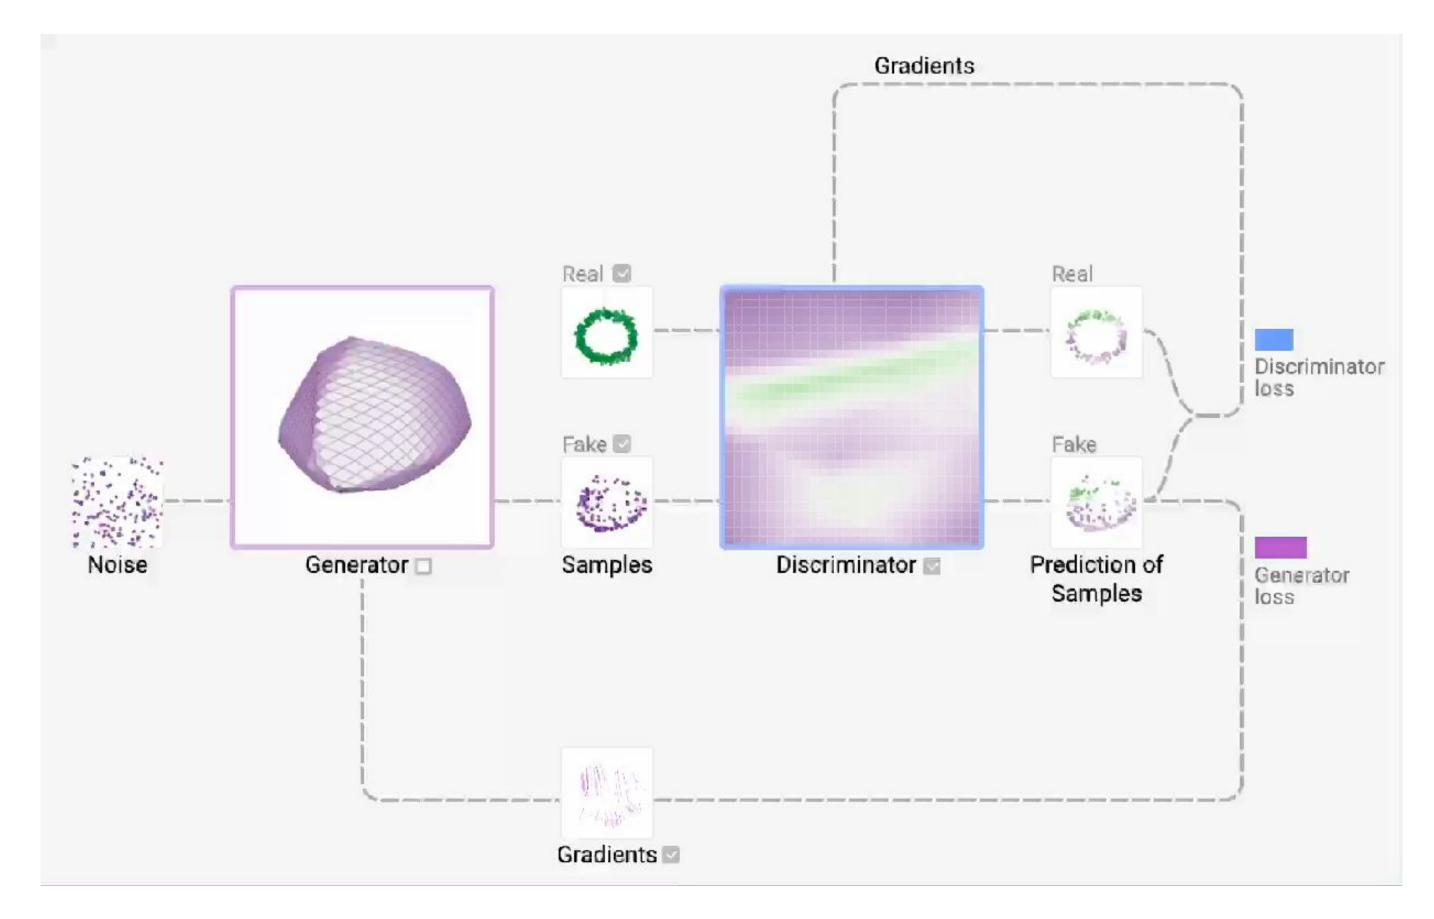
- Add to autoencoders the ability to generate new samples
  - Probabilistic latent space (assumed to be multivariate Gaussian)
- ullet Minimize at the same time reconstruction error and KL divergence between latent and  $N(0,{f I})$
- Sampling from the latent distribution allows generation and interpolation

#### GENERATIVE ADVERSARIAL NETWORKS Recap

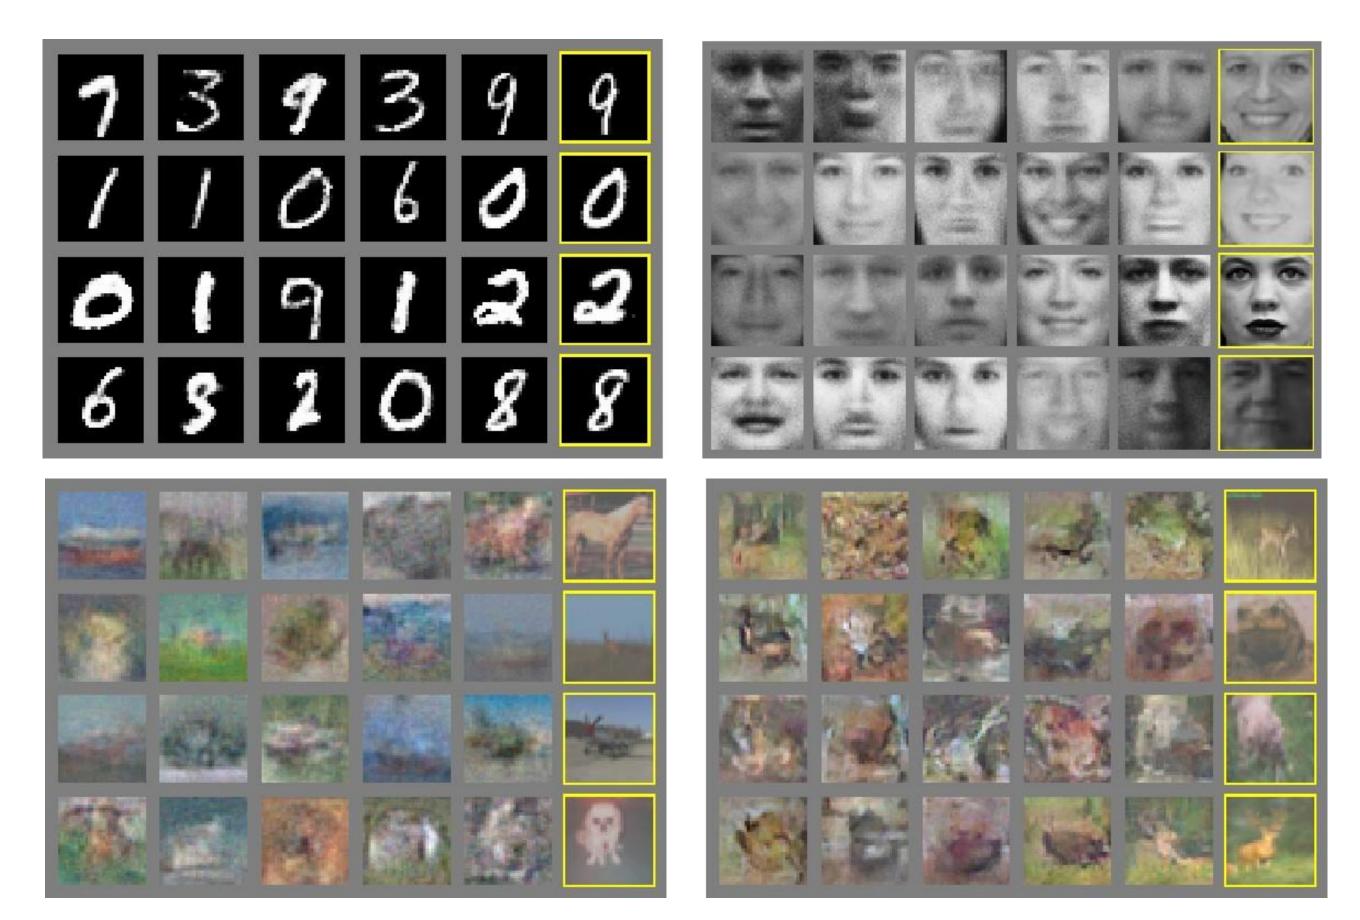


- Generator starts from Gaussian noise and generates a data point in order to fool the discriminator
- Discriminator tries to distinguish between real and fake data
- Training becomes a games between generator and discriminator
  - Solution lies in the Nash equilibrium between the two participants

## GAN VISUALIZATION



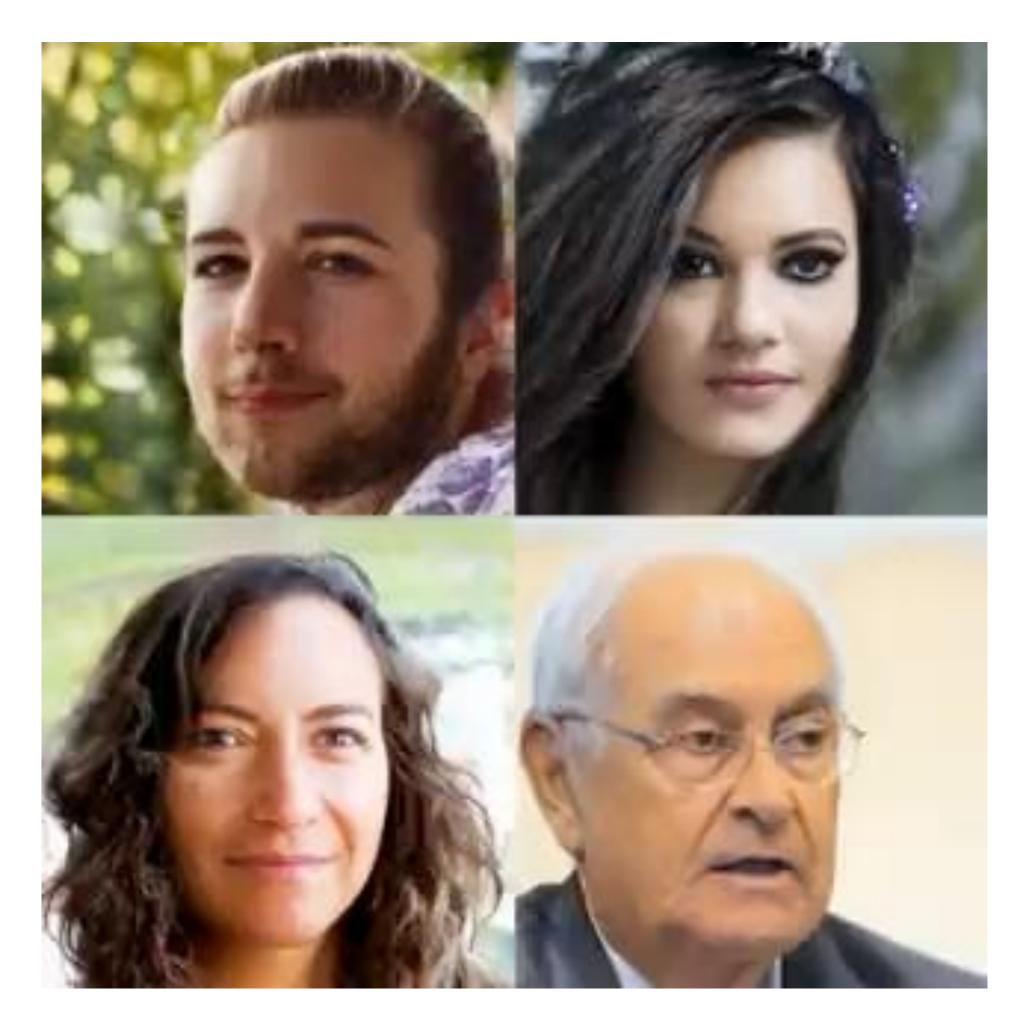
#### RESULTS - 2016

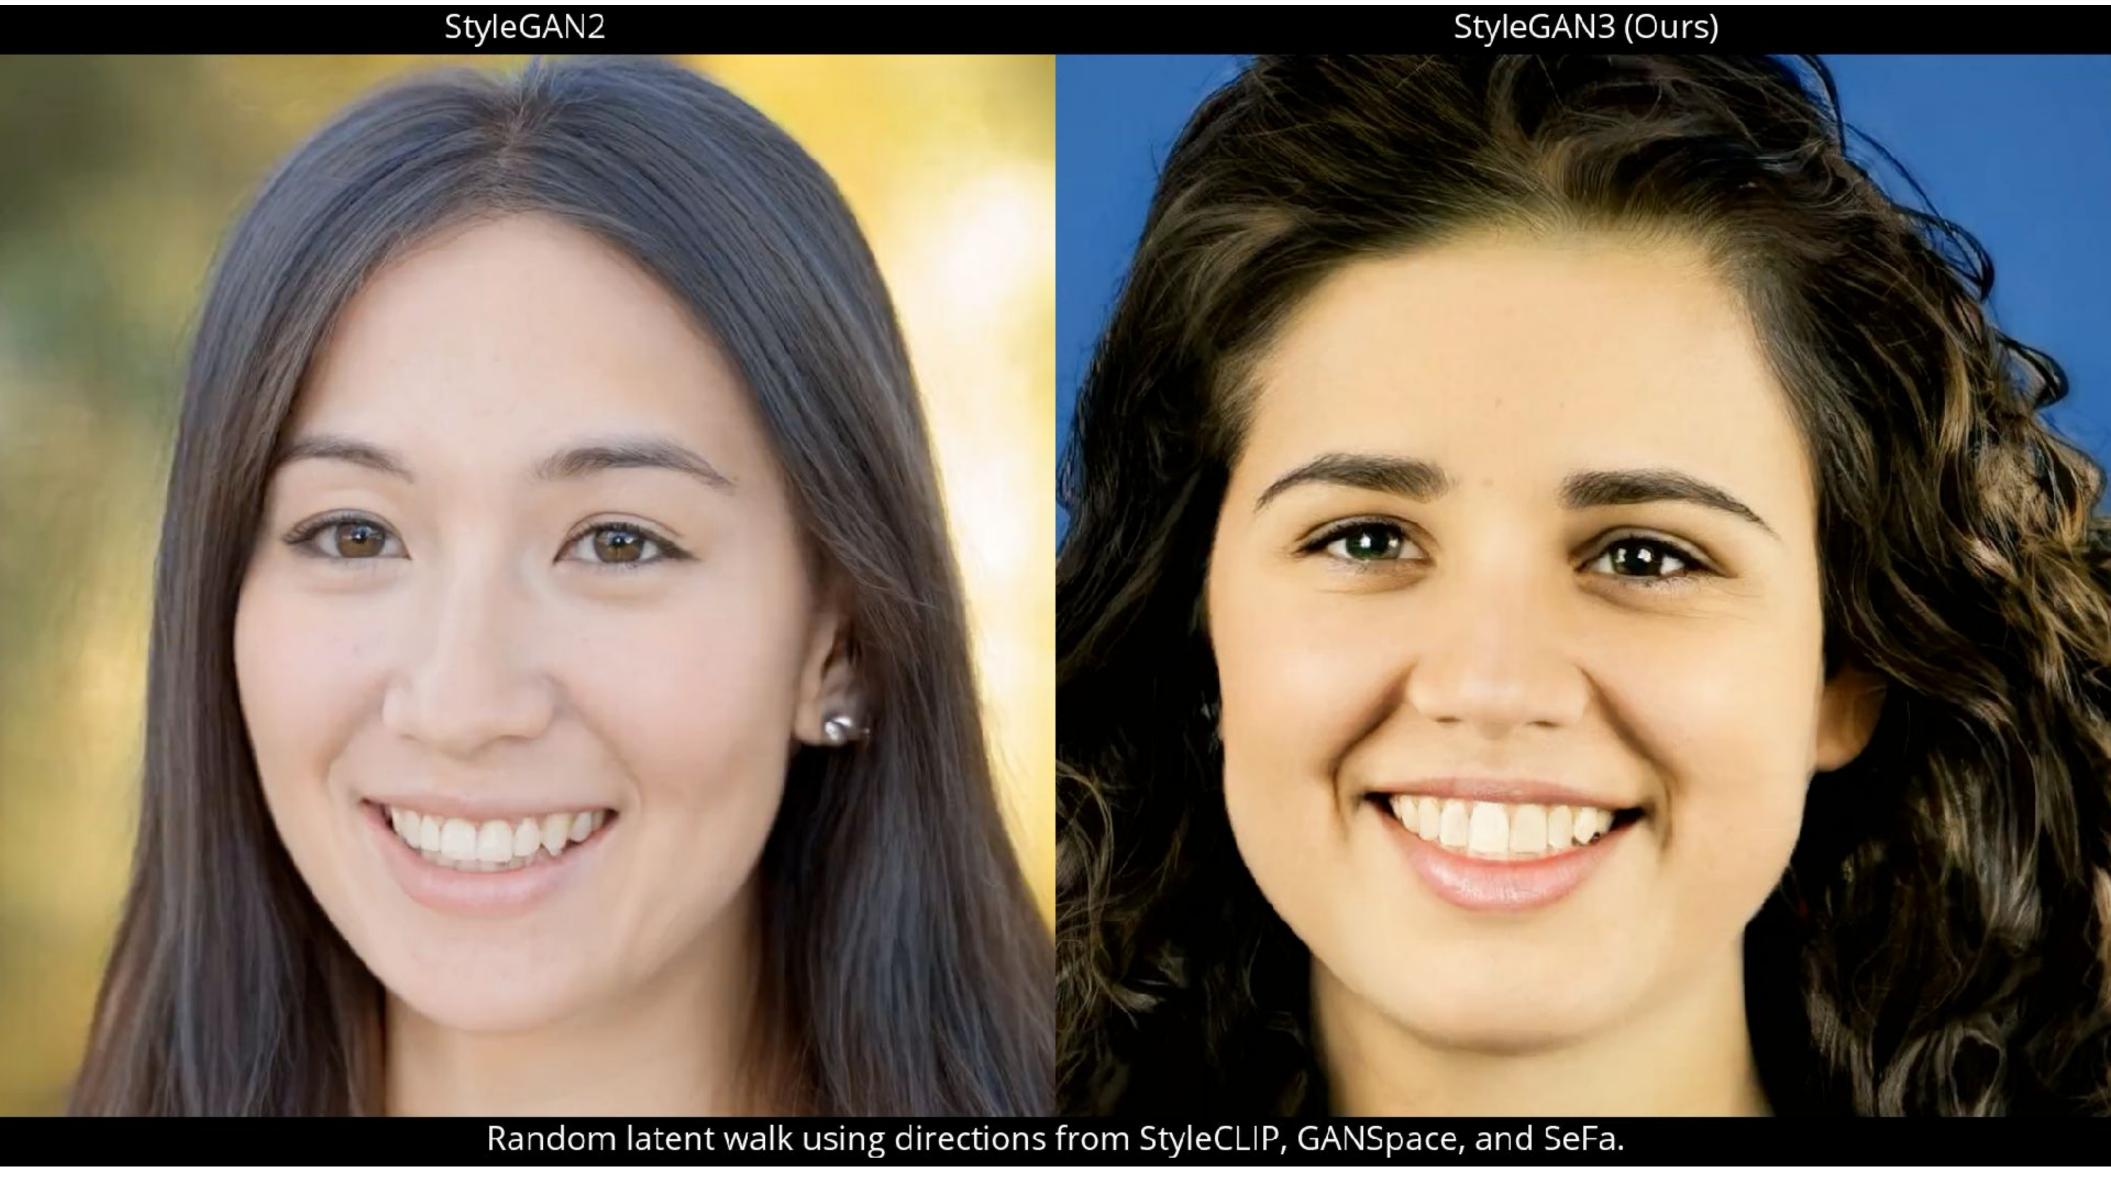


MNIST, Toronto Face Dataset, FC CIFAR-10, Con CIFAR-10 results

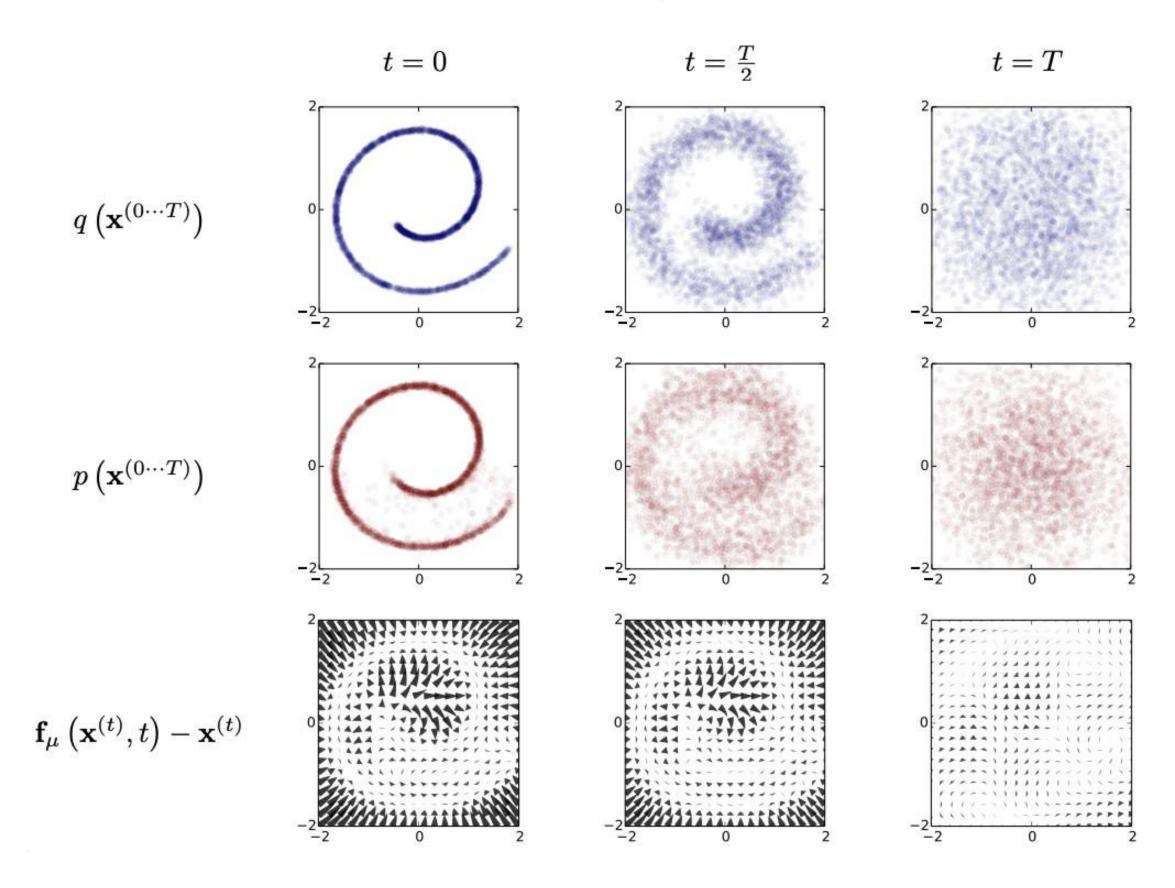
#### RESULTS - 2021

- Huge improvements thanks to several tricks, especially in the StyleGAN family
- Smooth walking in the latent space bears meaningful change in the output





### DEEP UNSUPERVISED LEARNING USING NONEQUILIBRIUM THERMODYNAMICS



Sohl-Dickstein, J., Weiss, E., Maheswaranathan, N., & Ganguli, S. (2015, June). Deep unsupervised learning using nonequilibrium thermodynamics. In International Conference on Machine Learning (pp. 2256-2265). PMLR.

#### DENOISING DIFFUSION PROBABILISTIC MODEL

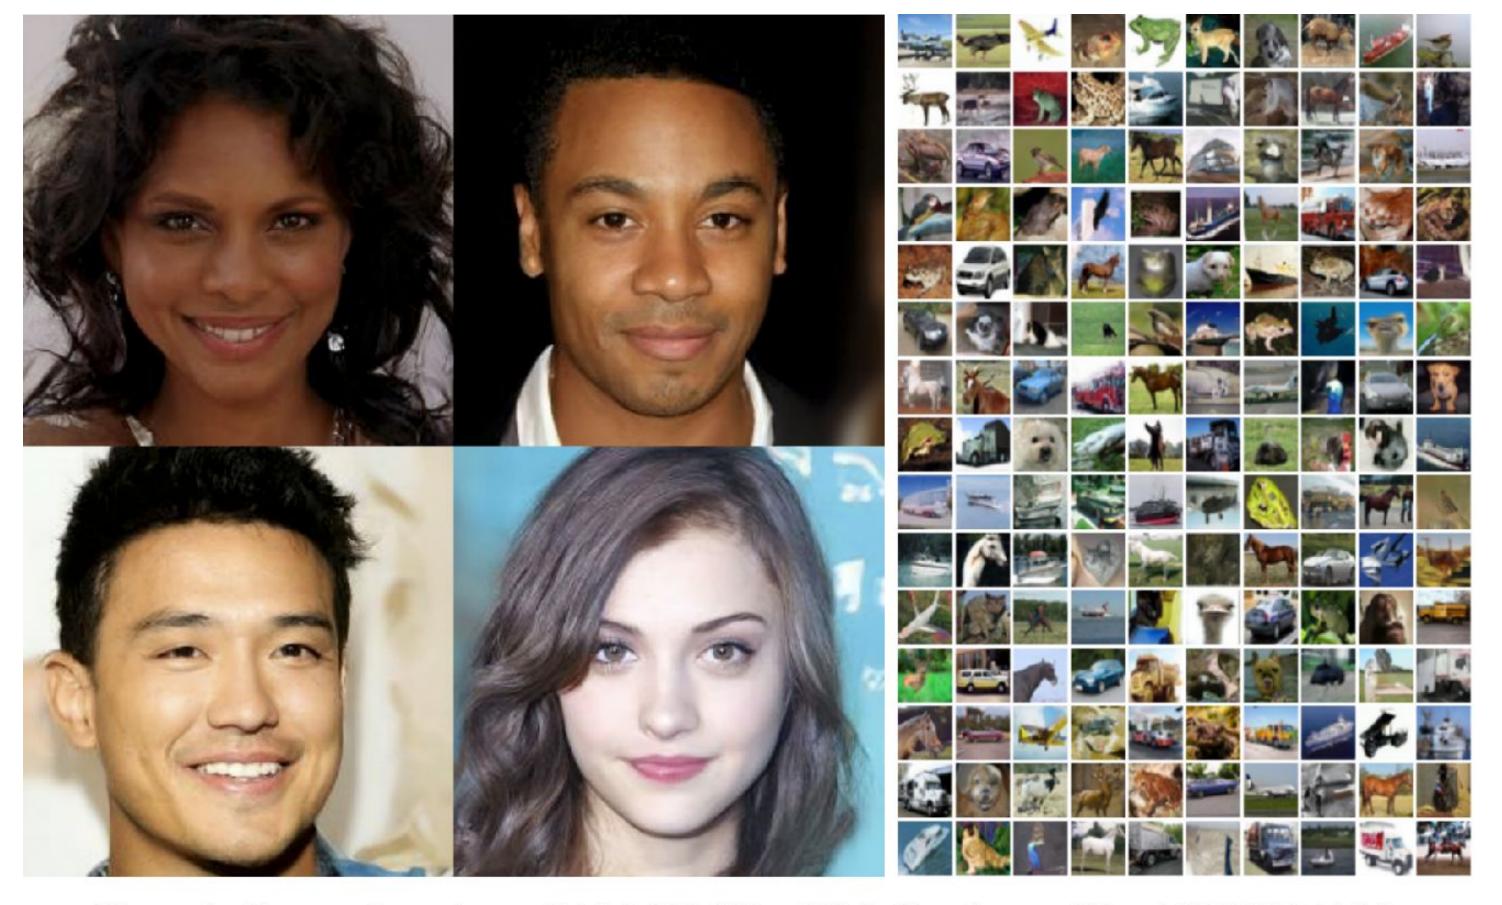


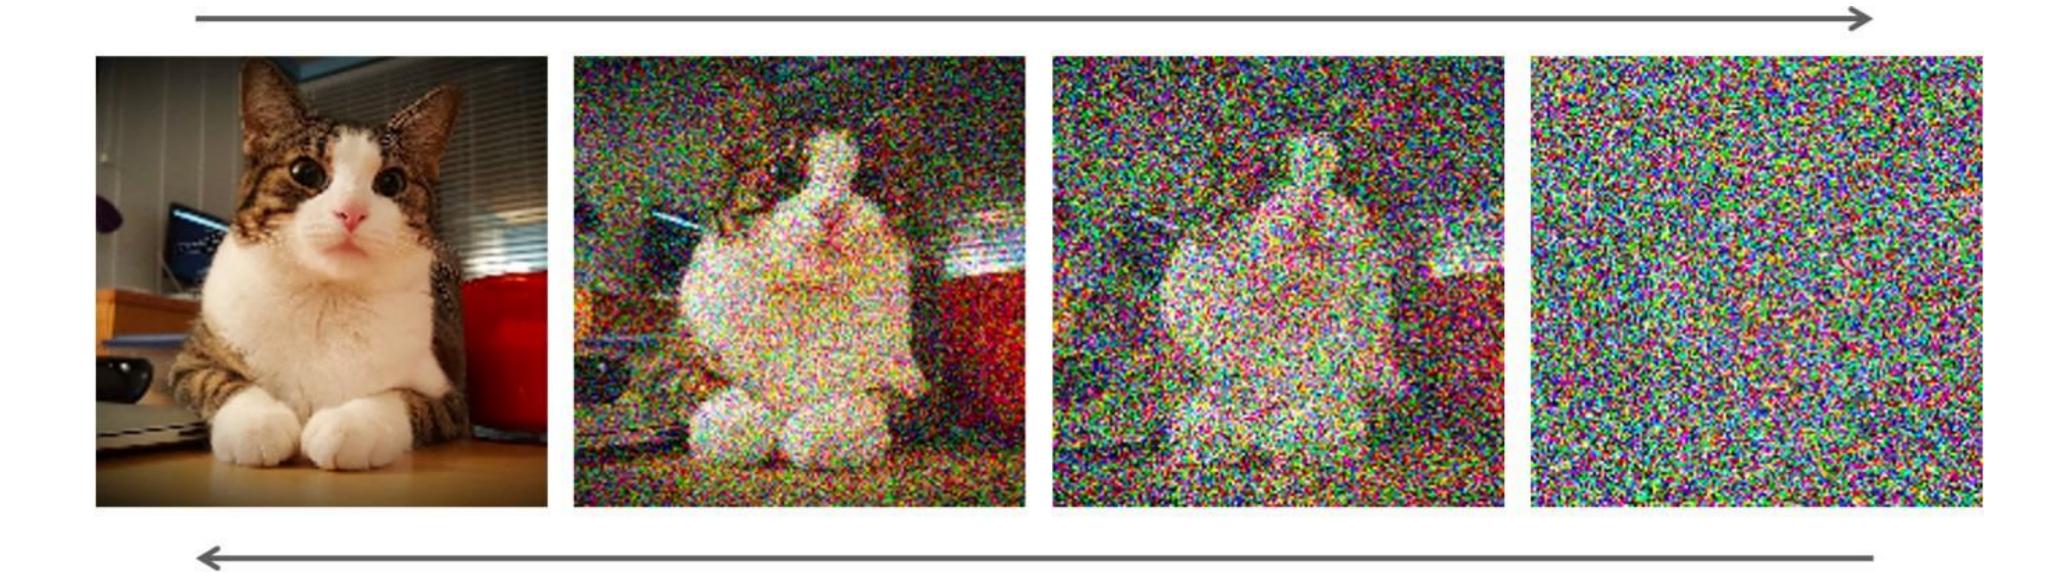
Figure 1: Generated samples on CelebA-HQ 256 × 256 (left) and unconditional CIFAR10 (right)

Ho, J., Jain, A., & Abbeel, P. (2020). Denoising diffusion probabilistic models. Advances in Neural Information Processing Systems, 33, 6840-6851.

# THEORY

È importante che imparimo il flavour e l'algoritmo

#### BASIC FORMULATION

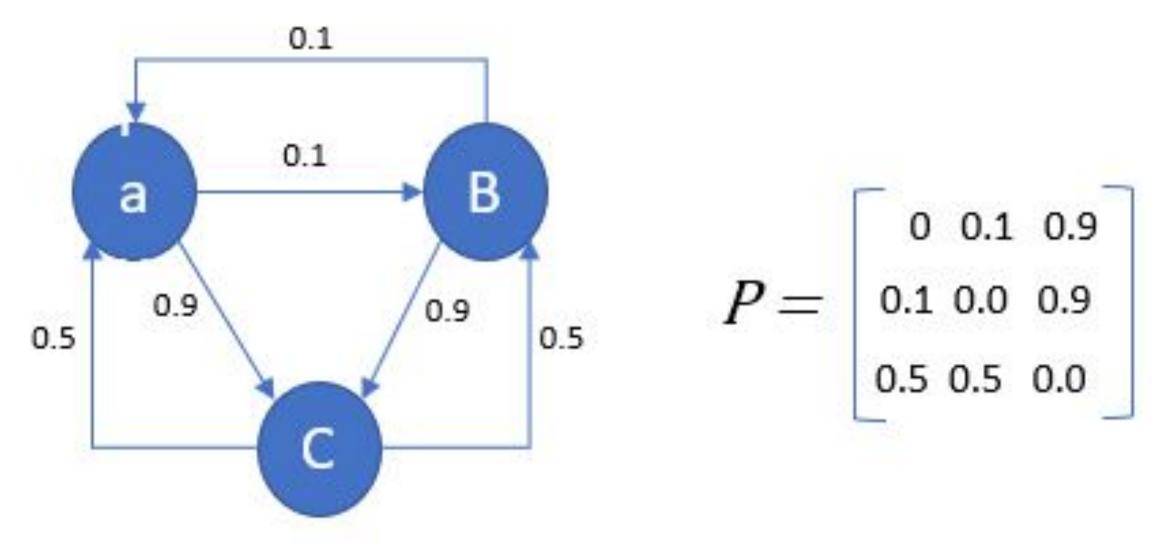


Denoising diffusion models consists of two processes:

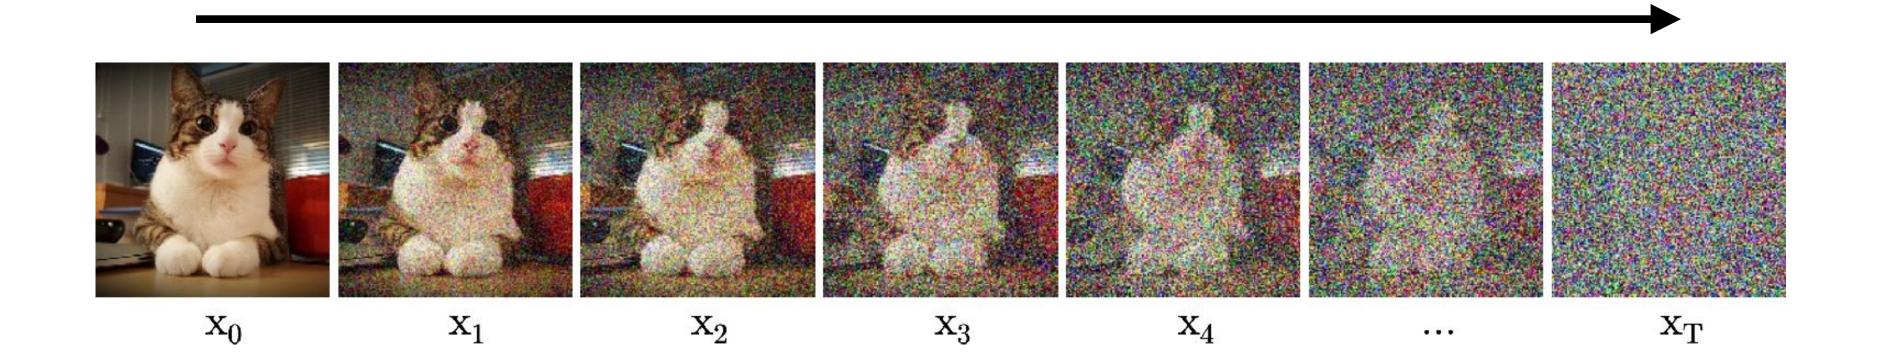
- Forward process to add noise
- Reverse process denoises to generate data

#### RECAP: MARKOV CHAIN

A Markov chain or Markov process is a stochastic model describing a sequence of possible events in which the probability of each event depends only on the state attained in the previous event

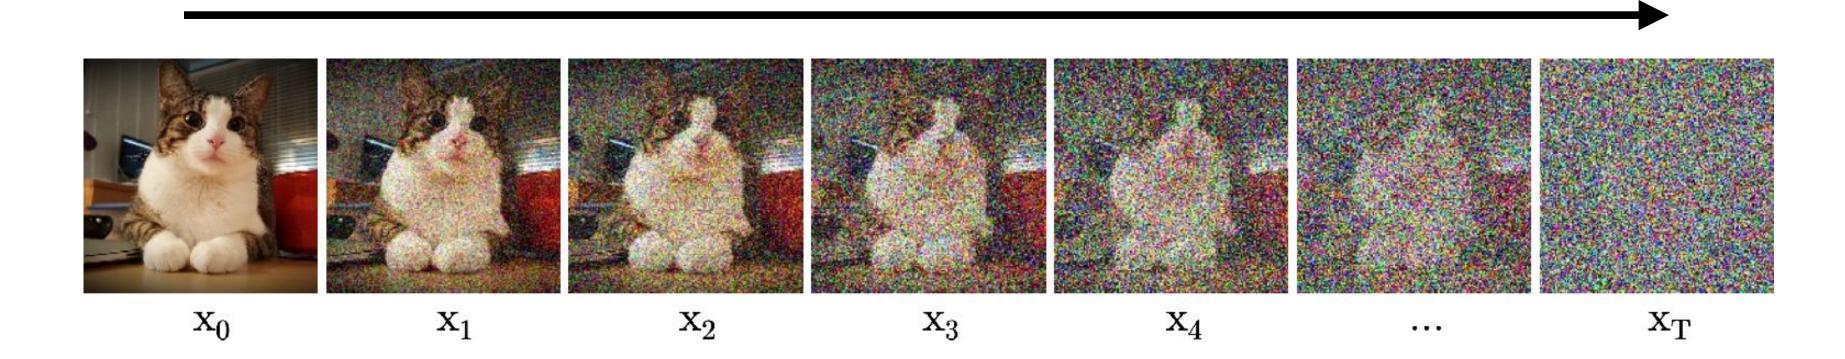


gli archi rappresentano la probabilità di muoversi da un evento all'altro



Formally we have the following formulation:

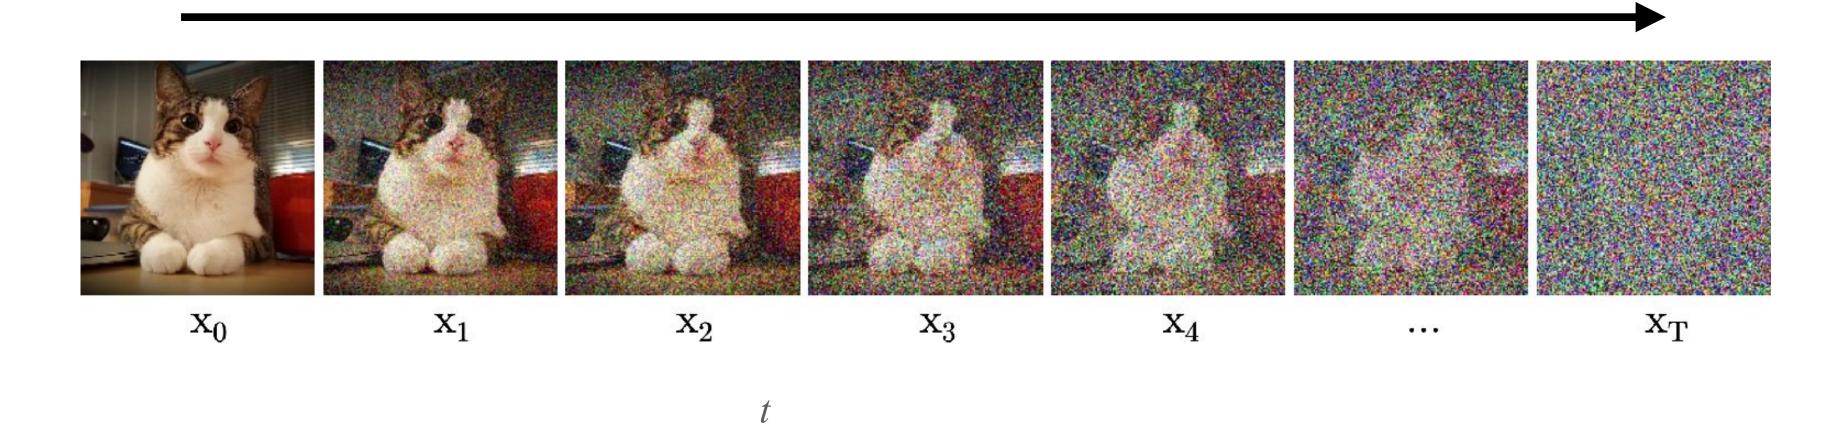
$$q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{1-\beta_t}x_{t-1}, \beta_t \mathbf{I}) \implies q(x_{1:T}|x_0) = \prod_{t=1}^{T} q(x_t|x_{t-1})$$



$$q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{1-\beta_t}x_{t-1}, \beta_t \mathbf{I}) \implies q(x_{1:T}|x_0) = \prod_{t=1}^{T} q(x_t|x_{t-1})$$

With a change of variables:

Define: 
$$\alpha_t = (1 - \beta_t) \implies \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s) \implies q(x_t | x_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

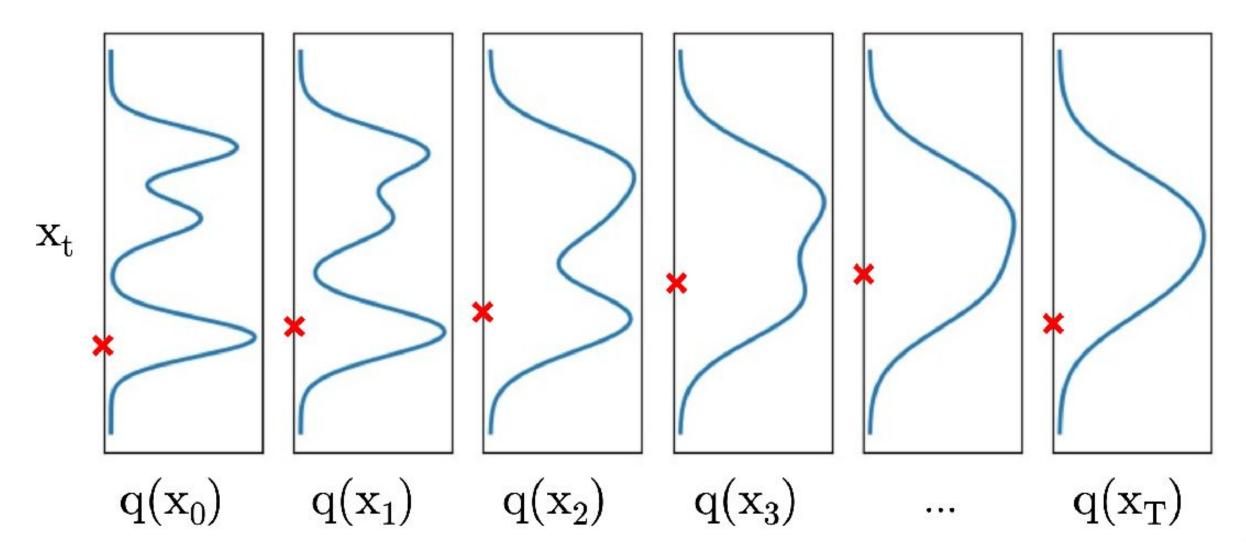


We can then sample directly at the desired timestep  $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{(1-\bar{\alpha}_t)} \epsilon, \epsilon \sim \mathcal{N}(0,\mathbf{I})$ 

 $\beta_t$  is the noise schedule

- Formally we are applying a
   Gaussian convolution to the data at each timestep
- Practically we are smoothening out the distribution to a Gaussian one

#### **Diffused Data Distributions**

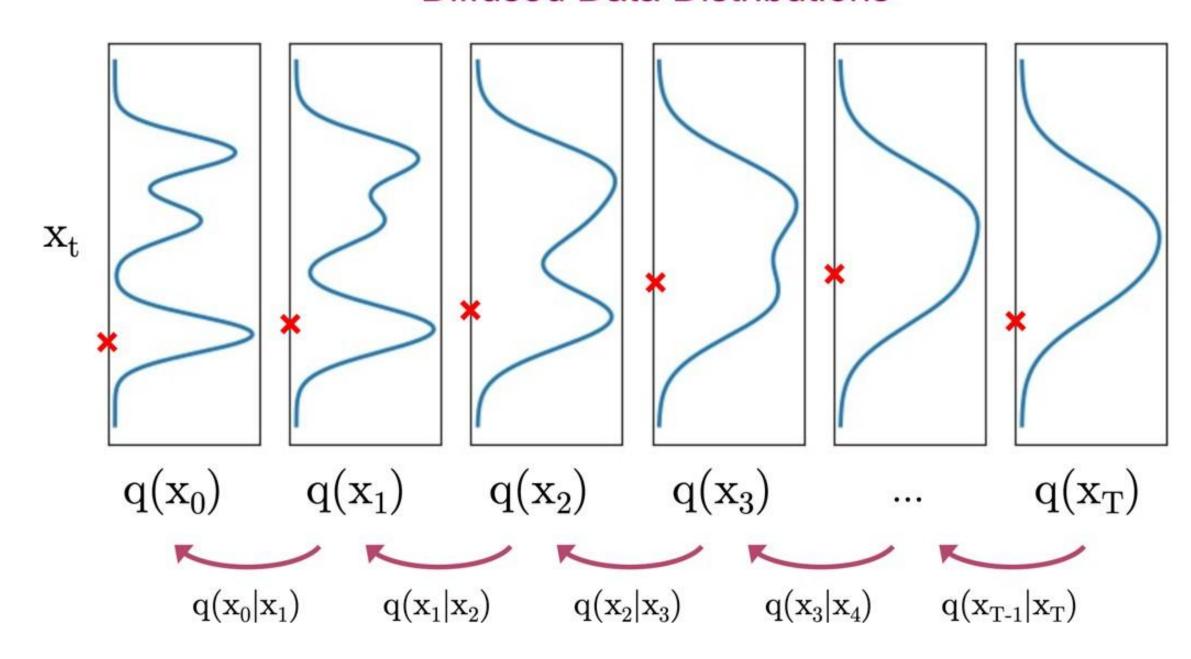


#### GENERATION PROCESS

#### Given that $q(x_T) \approx \mathcal{N}(0, \mathbf{I})$ :

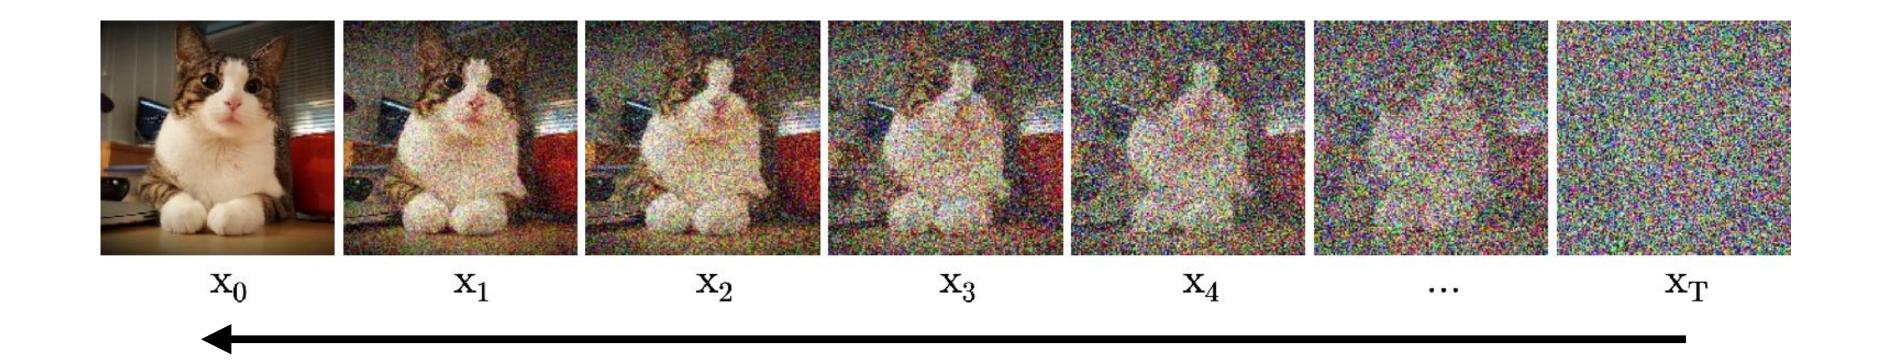
- Sample  $x_T \sim \mathcal{N}(0, \mathbf{I})$
- Iteratively sample  $x_{t-1} = q(x_{t-1} | x_t)$

#### **Diffused Data Distributions**



$$q(x_{t-1} | x_t) \propto q(x_{t-1})q(x_t | x_{t-1})$$
 is generally **intractable**, but we can **approximate** with another Gaussian

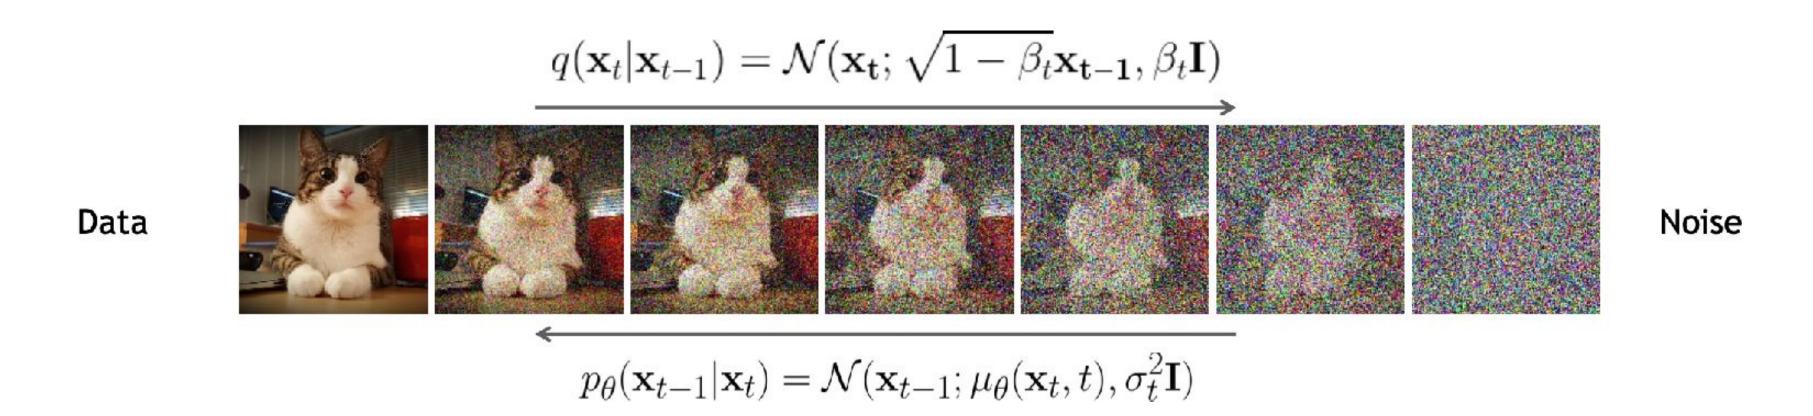
#### REVERSE PROCESS



$$p(x_T) \sim \mathcal{N}(0, \mathbf{I})$$

$$p_{\theta}(x_{t-1} \mid x_t) \sim \mathcal{N}\left[\mu_{\theta}(x_t, t) \mid \sigma_t^2 \mathbf{I}\right]$$
Trainable network
$$p_{\theta}(x_{t-1} \mid x_t) \sim \mathcal{N}\left[\mu_{\theta}(x_t, t) \mid \sigma_t^2 \mathbf{I}\right]$$

#### NOISING SCHEDULE



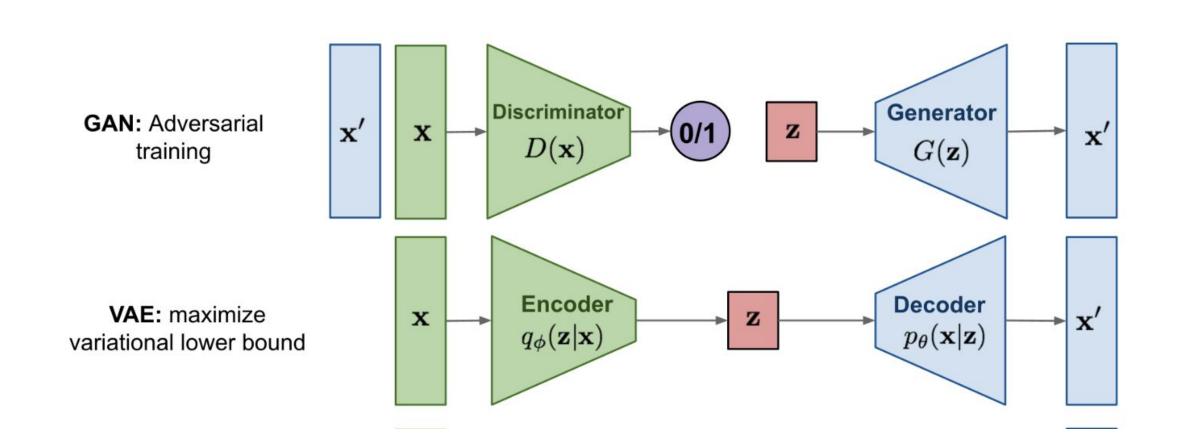
We can control the variance of the forward diffusion and reverse denoising processes respectively. Often a linear schedule is used for  $\beta_t$ , and  $\sigma_t^2$  is set equal to  $\beta_t$ .

Kingma et al. [NeurIPS 2022] introduce a new parameterization of diffusion models using signal-to-noise ratio (SNR), and show how to learn the noise schedule by minimizing the variance of the training objective.

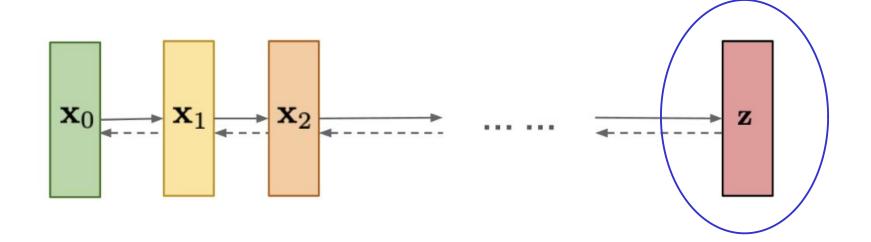
Other improvements: (Improved DPM by Nichol and Dhariwal [ICML 2021], Analytic-DPM by Bao et al. [ICLR 2022]).

#### CONNECTION TO VAE, GANS

- Latent variables have the same dimensionality of data
- The same model is applied across different timesteps
- The model is trained by reweighing the variational bound



#### Diffusion models: Gradually add Gaussian noise and then reverse



#### TRAINING PARAMETRISATION

We can train the model in a similar fashion as VAE, with a Variational Upper Bound

$$L = \mathbb{E}_{q(x_0)} \left[ -\log p_{\theta}(x_0) \right] \le \mathbb{E}_{q(x_0)q(x_{1:T}|x_0)} \left[ -\log \frac{p_{\theta}(x_{0:T})}{q(x_{1:T}|x_0)} \right]$$

These can be divided into three terms

$$L = \mathbb{E}_{q} \left[ D_{KL}(q(x_{T}|x_{0}) \mid |p(x_{T})) + \sum_{t>1} D_{KL}(q(x_{t-1}|x_{t},x_{0}) \mid |p_{\theta}(x_{t-1}|x_{t})) - \log p_{\theta}(x_{0}|x_{1}) \right]$$

This is the original loss. Can this be simplified?

#### TRAINING PARAMETRISATION

$$L = \mathbb{E}_{q} \left[ D_{KL}(q(x_{T} | x_{0}) | | p(x_{T})) + \sum_{t>1} D_{KL}(q(x_{t-1} | x_{t}, x_{0}) | | p_{\theta}(x_{t-1} | x_{t})) - \log p_{\theta}(x_{0} | x_{1}) \right]$$

KL between Gaussians has a nice closed form, but Ho (with some math) proves the training can be simplified to a noise prediction problem

Skipping some math steps, we obtain a new loss

$$L_{simple} = \mathbb{E}_{x_0 \sim q(x_0), \epsilon \sim \mathcal{N}(0, \mathbf{I}), t \sim \mathcal{U}(1, T)} \left[ \left| \left| \epsilon - \epsilon_{\theta} (\sqrt{\bar{\alpha}_t} x_0 + \sqrt{(1 - \bar{\alpha}_t)} \epsilon, t) \right| \right|^2 \right]$$

#### TRAINING AND SAMPLING

In questo caso il training è più semplice dell'inference phase

# Algorithm 1 Training 1: repeat 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 3: $t \sim \text{Uniform}(\{1, \dots, T\})$

- 4:  $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on

$$\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta} \left( \sqrt{\bar{\alpha}_{t}} \mathbf{x}_{0} + \sqrt{1 - \bar{\alpha}_{t}} \boldsymbol{\epsilon}, t \right) \right\|^{2}$$

6: until converged

#### **Algorithm 2** Sampling

1: 
$$\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

2: **for** 
$$t = T, ..., 1$$
 **do**

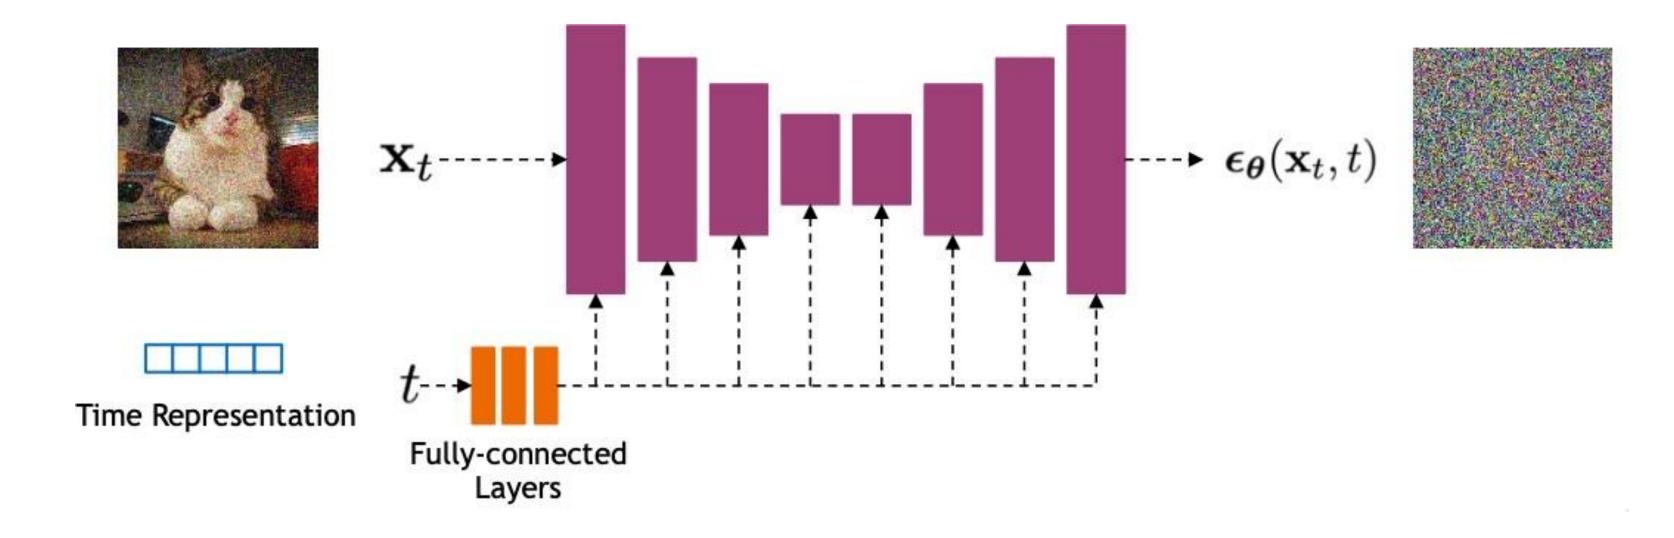
3: 
$$\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

4: 
$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$$

5: end for

6: return  $x_0$ 

#### NETWORK

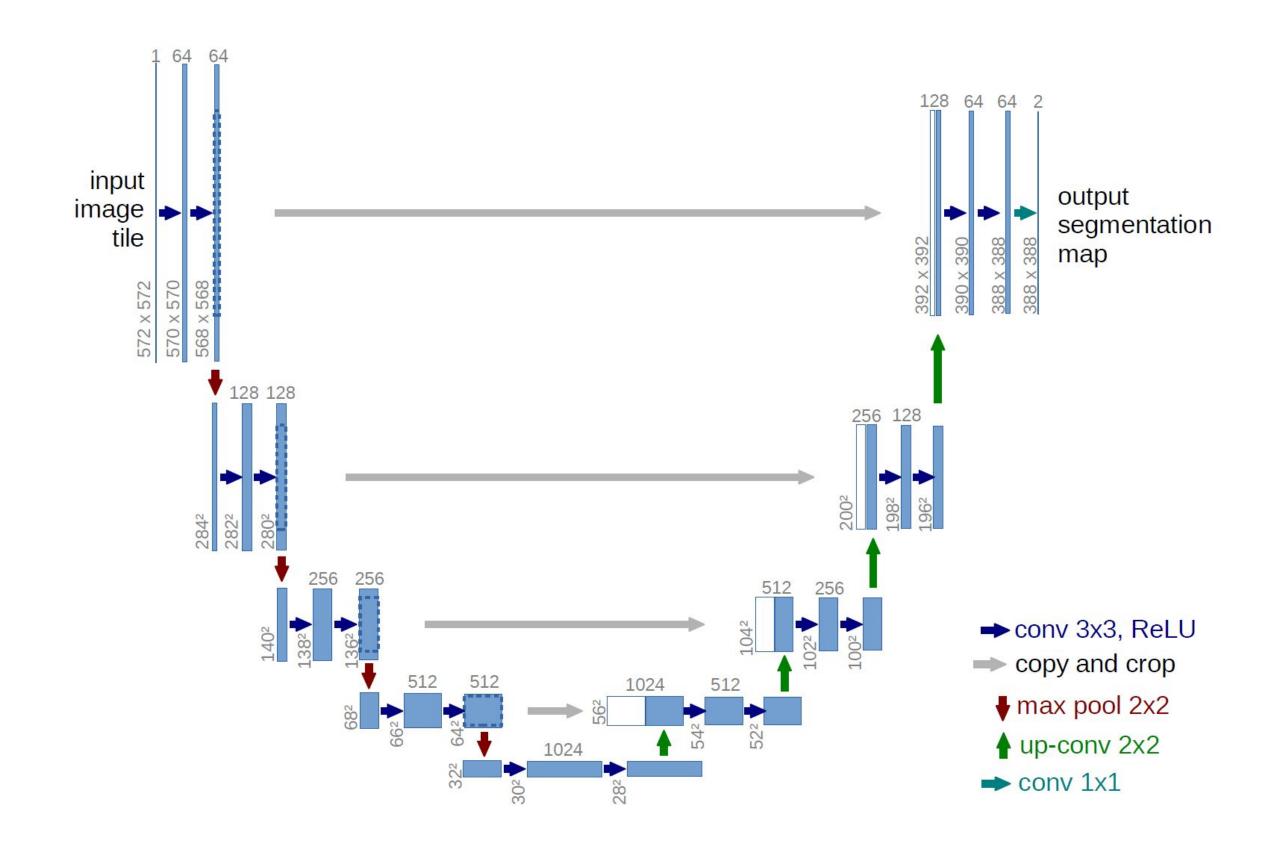


Pretty much free choice on the network architecture, no theoretical constraints on this For images use **U-Net** (with attention)

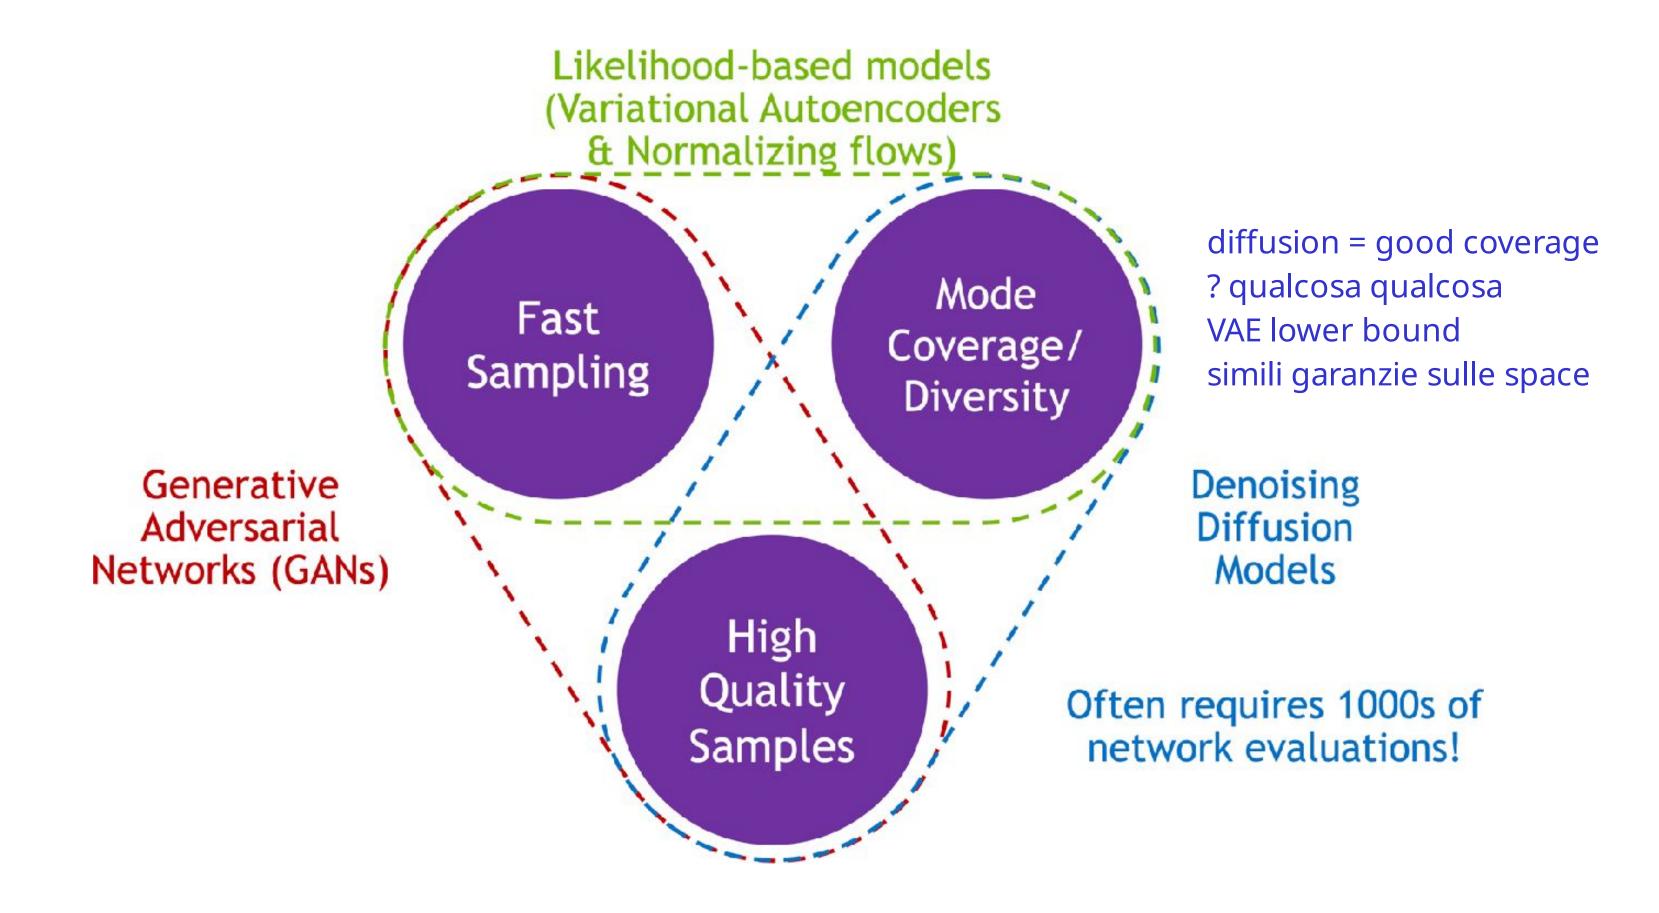
Time features are usually **sinusoidal** of random **Fourier** features. How to embed them in the network is another free choice (e.g. spatial concatenation, AdalN, etc)

#### U-NET

- The U-NET architecture contains two paths.
- First path is the contraction path (also called as the encoder) which is used to capture the context in the image. The encoder is just a traditional stack of convolutional and max pooling layers.
- The second path is the symmetric expanding path (also called as the decoder) which is used to enable precise localization using transposed convolutions.
- It is an end-to-end fully convolutional network (FCN), i.e. it only contains
   Convolutional layers and does not contain any Dense layer because of which it can accept image of any size.



#### GENERATIVE TRILEMMA

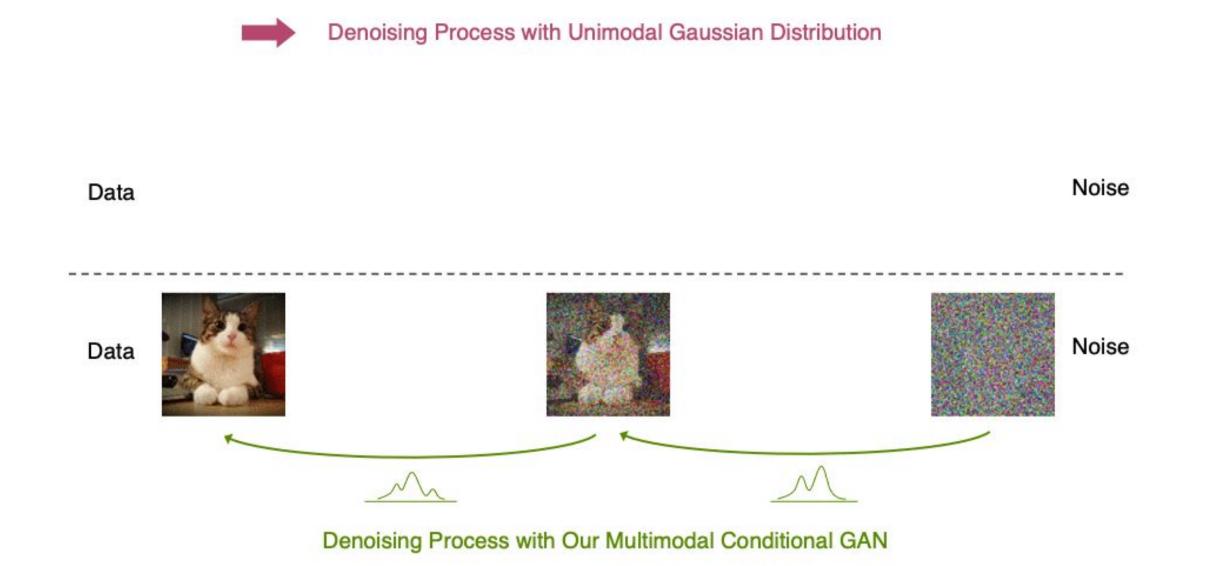


## ADVANCED TRICKS

Il resto è extra

#### DIFFUSION GANS

Generative denoising diffusion models typically assume that the denoising distribution can be modeled by a Gaussian distribution. This assumption holds only for small denoising steps, which in practice translates to thousands of denoising steps in the synthesis process. In diffusion GANs, the denoising model is represented using multimodal and complex conditional GANs, enabling to efficiently generate data in a few steps.

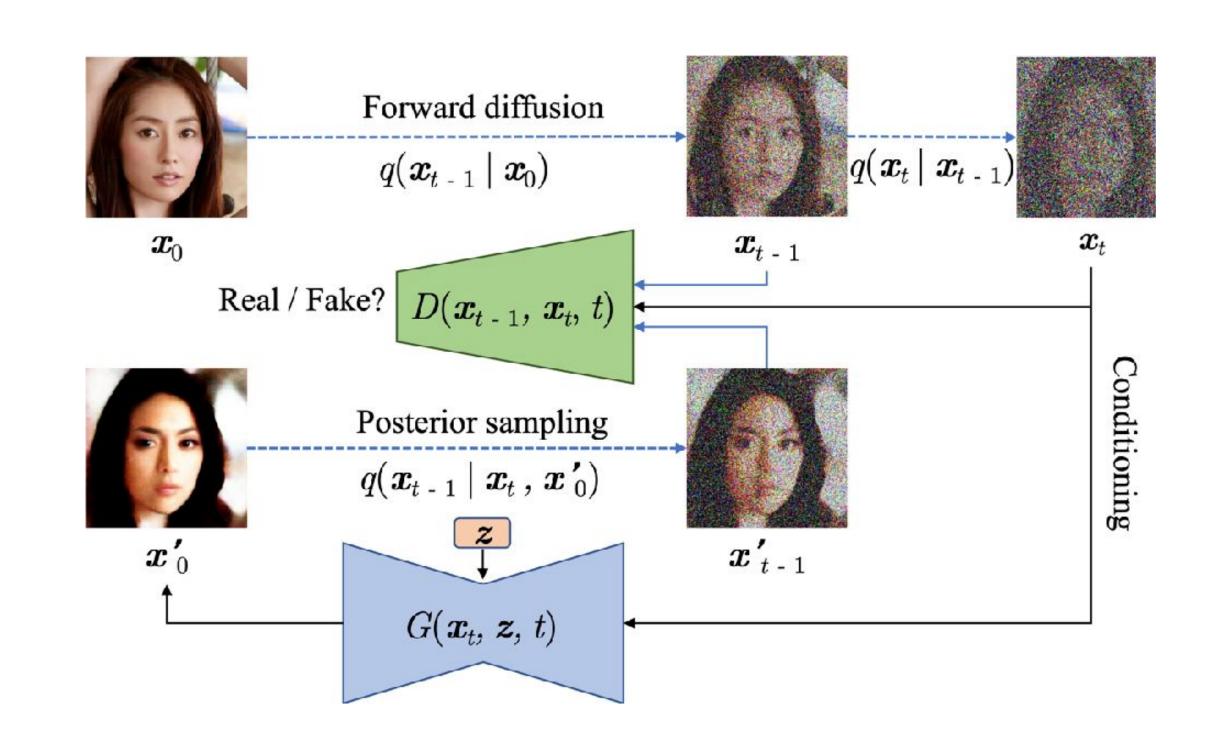


Xiao, Z., Kreis, K., & Vahdat, A. (2021). Tackling the generative learning trilemma with denoising diffusion gans. ICLR 2022.

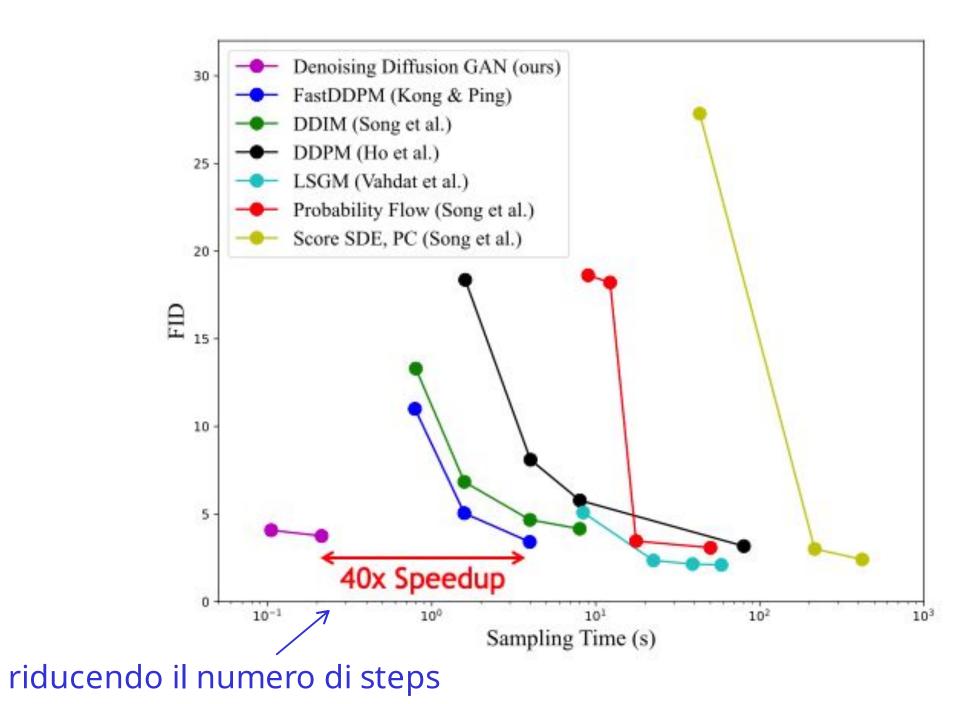
#### DIFFUSION GANS

Compared to a one-shot GAN generator:

- Both generator and discriminator are solving a much simpler problem.
- Stronger mode coverage
- Better training stability



#### DIFFUSION GANS



Sample Quality vs. Sampling Time

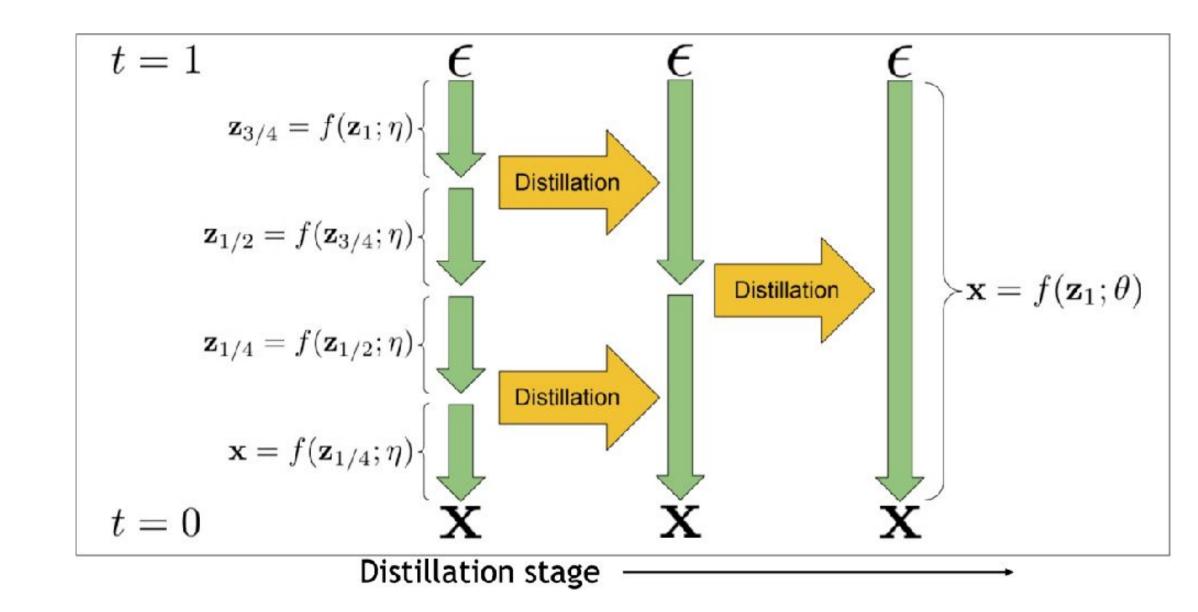


CIFAR-10 Samples

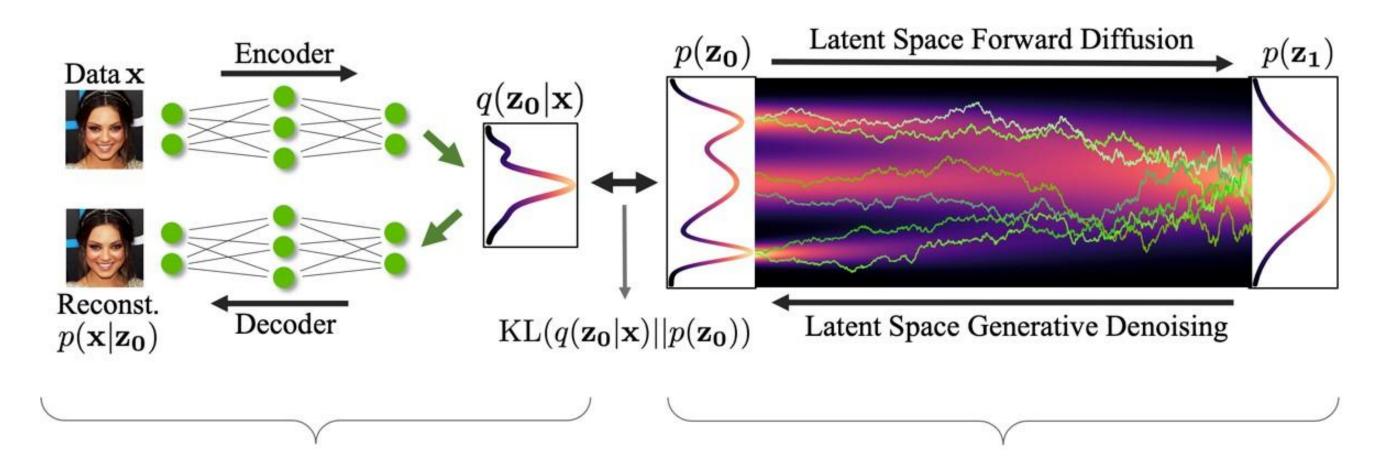
Xiao, Z., Kreis, K., & Vahdat, A. (2021). Tackling the generative learning trilemma with denoising diffusion gans. ICLR 2022.

#### DISTILLATION

- **Distill** a deterministic DDIM sampler to the same model architecture.
- At each stage, a "student" model is learned to distill two adjacent sampling steps of the "teacher" model to one sampling step.
- At next stage, the "student" model from previous stage will serve as the new "teacher" model.



#### LATENT-SPACE DIFFUSION MODELS



Variational Autoencoder

**Denoising Diffusion Prior** 

- The distribution of **latent embeddings** is close to **Normal** distribution Simpler denoising and faster synthesis
- Augmented latent space
- Tailored autoencoders (graphs, text, 3D data, etc.)

#### TEXT-TO-IMAGE



"a hedgehog using a calculator"



"a corgi wearing a red bowtie and a purple party hat"



"robots meditating in a vipassana retreat"

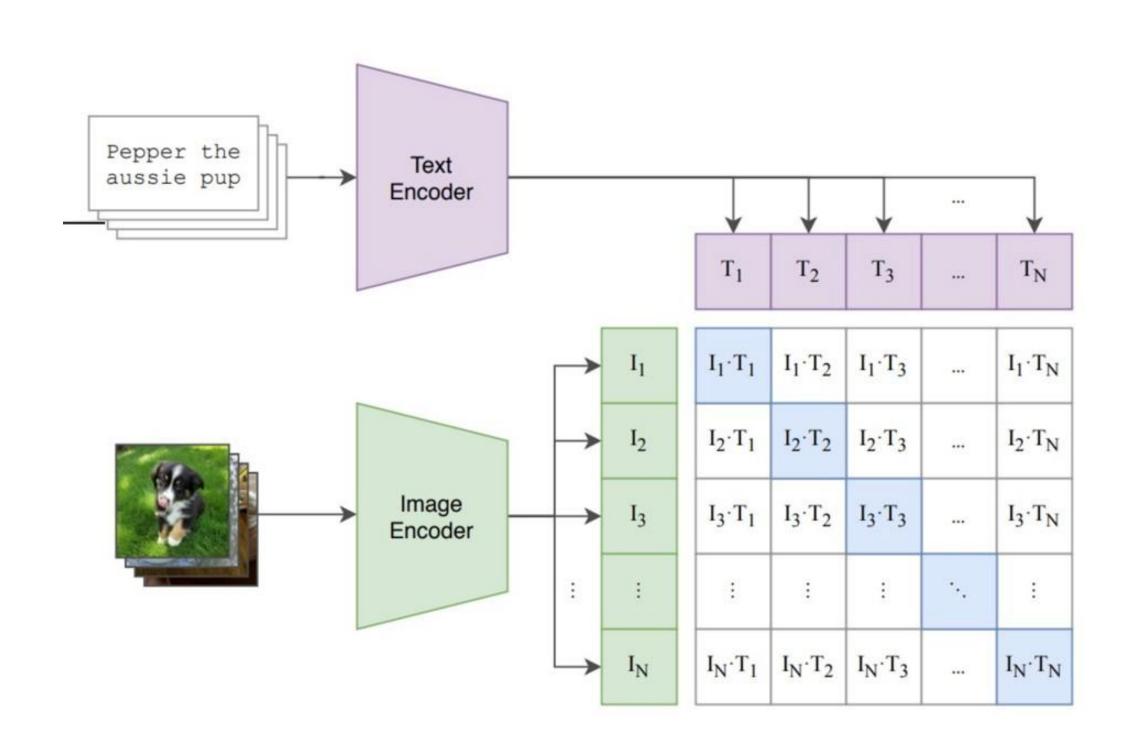


"a fall landscape with a small cottage next to a lake"

Nichol, A., Dhariwal, P., Ramesh, A., Shyam, P., Mishkin, P., McGrew, B., ... & Chen, M. (2021). Glide: Towards photorealistic image generation and editing with text-guided diffusion models. arXiv preprint arXiv:2112.10741.

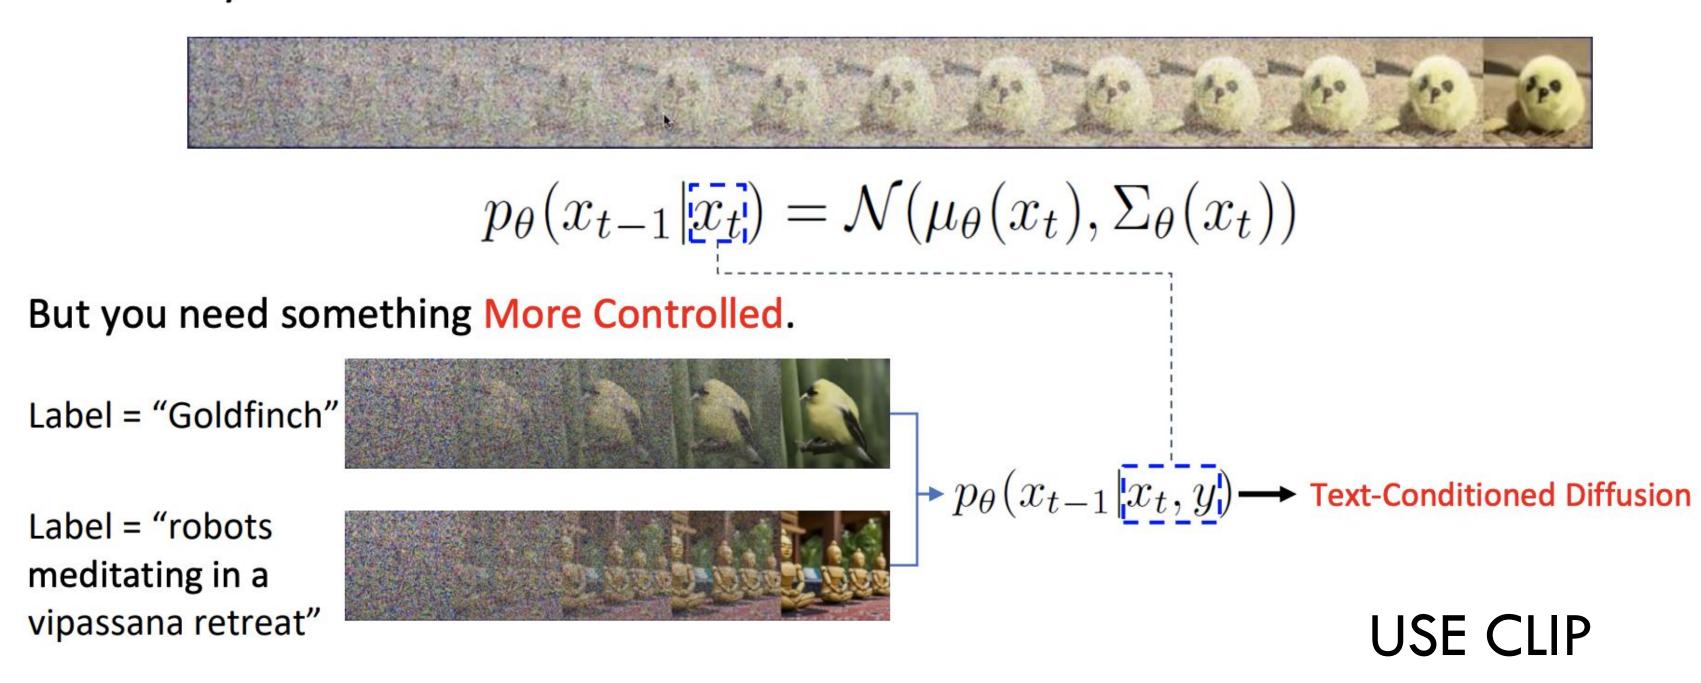
#### CONTRASTIVE LANGUAGE-IMAGE PERTAINING

- Jointly train a text encoder and an image encoder
- Train by maximising the similarity between embeddings of (text, image) pairs
- The resulting space has semantics for both images and text



#### GLIDE: IDEA

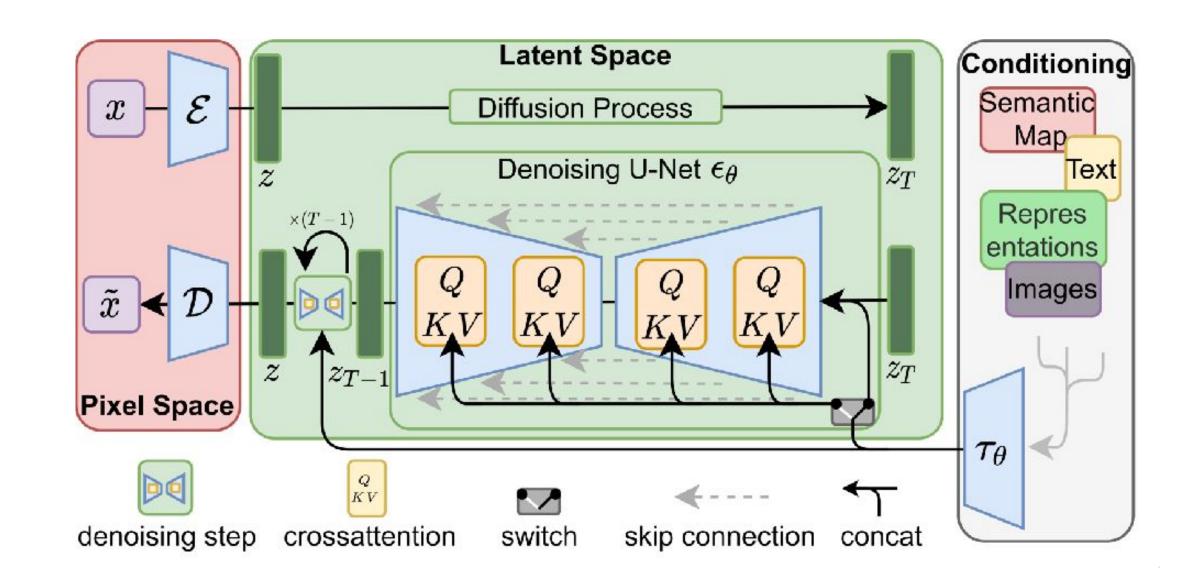
You already understand the Diffusion.



Nichol, A., Dhariwal, P., Ramesh, A., Shyam, P., Mishkin, P., McGrew, B., ... & Chen, M. (2021). Glide: Towards photorealistic image generation and editing with text-guided diffusion models. arXiv preprint arXiv:2112.10741.

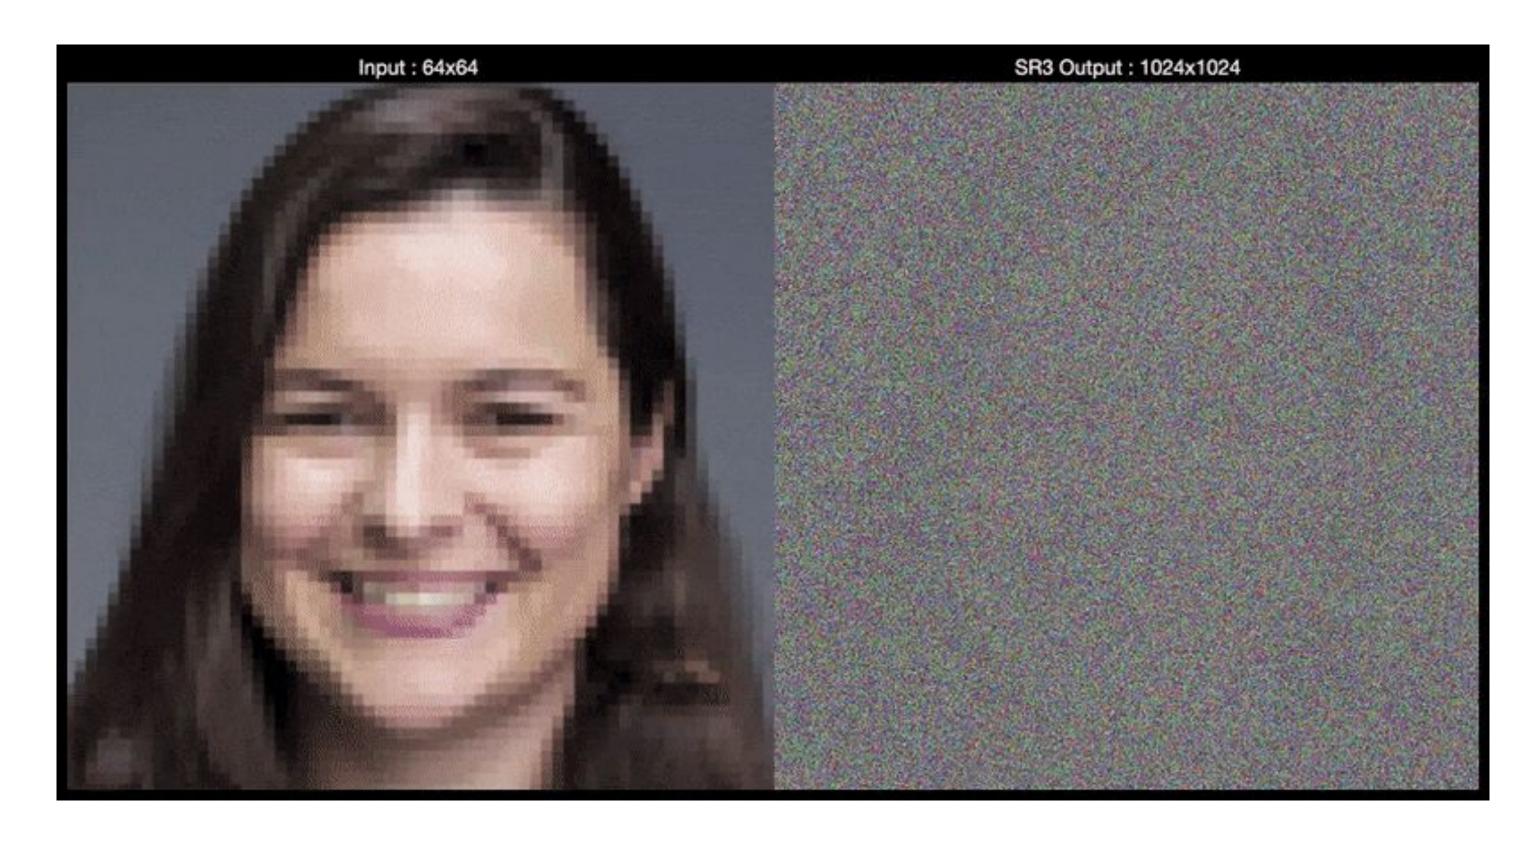
## TEXT-TO-IMAGE (STABLE DIFFUSION)

- Use CLIP embeddings as conditioning on latent diffusion via cross-attention
- Trained on a subset of LAION-5B (original dataset has 5 billion textimage pairs)
- Fast sampling thanks to diffusion in latent space



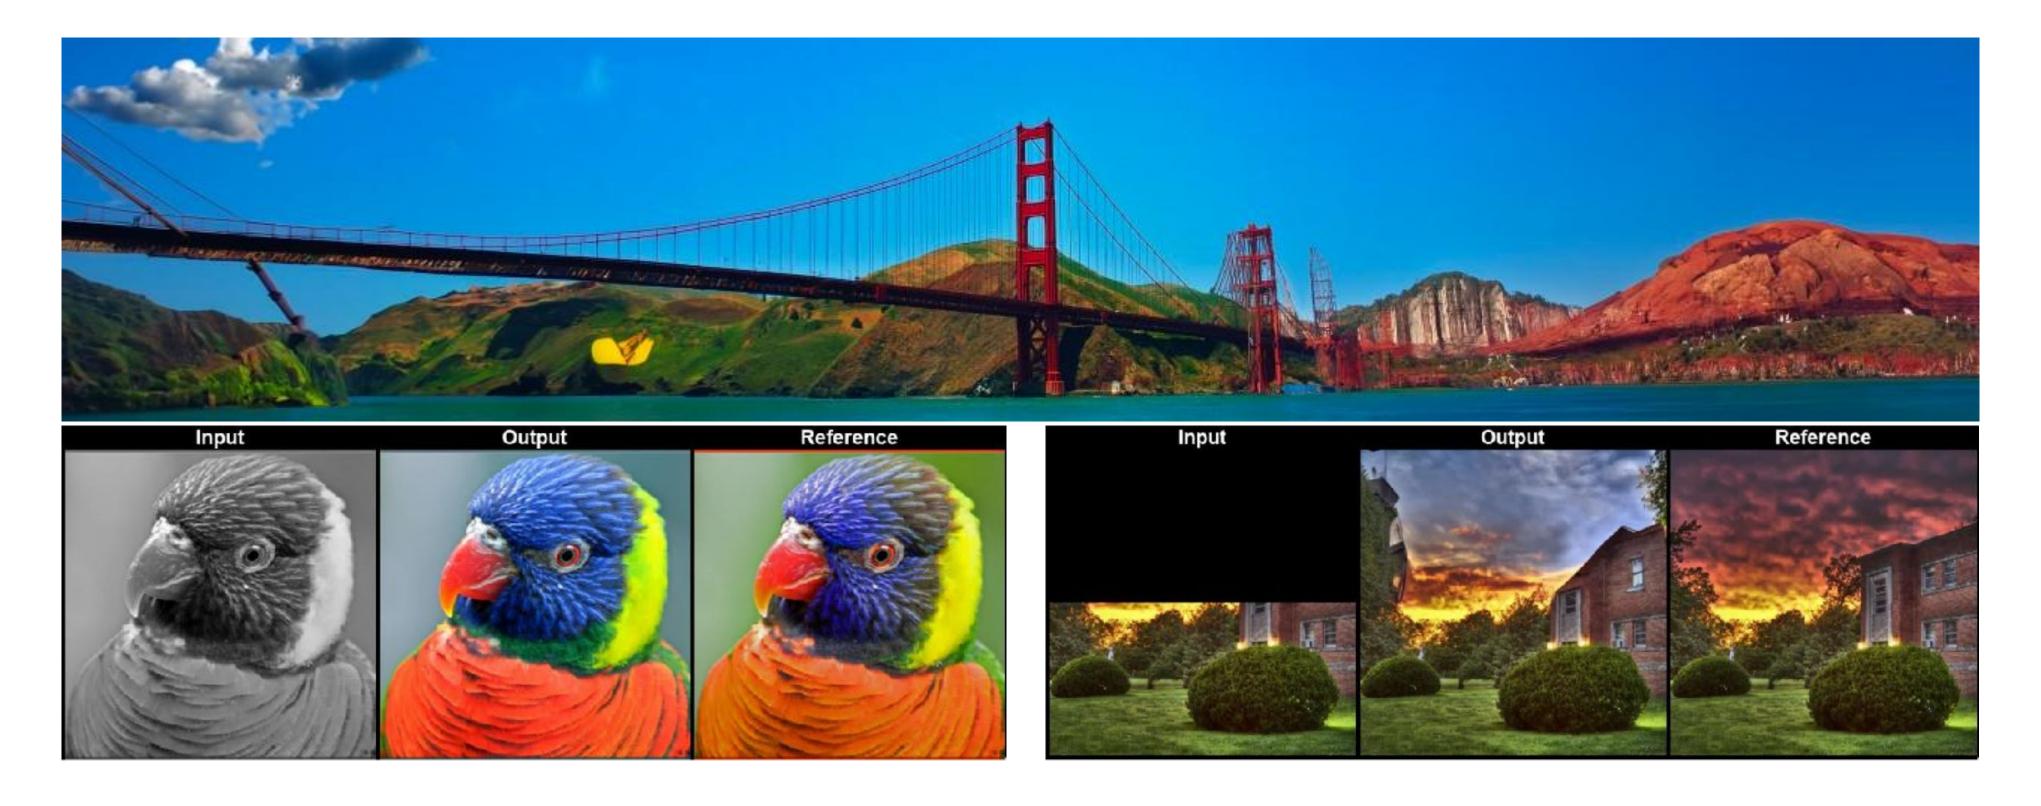
# DIFFUSION 200

#### SUPER RESOLUTION



Saharia, C., Ho, J., Chan, W., Salimans, T., Fleet, D. J., & Norouzi, M. (2022). Image super-resolution via iterative refinement. IEEE Transactions on Pattern Analysis and Machine Intelligence.

### IMAGE-TO-IMAGE



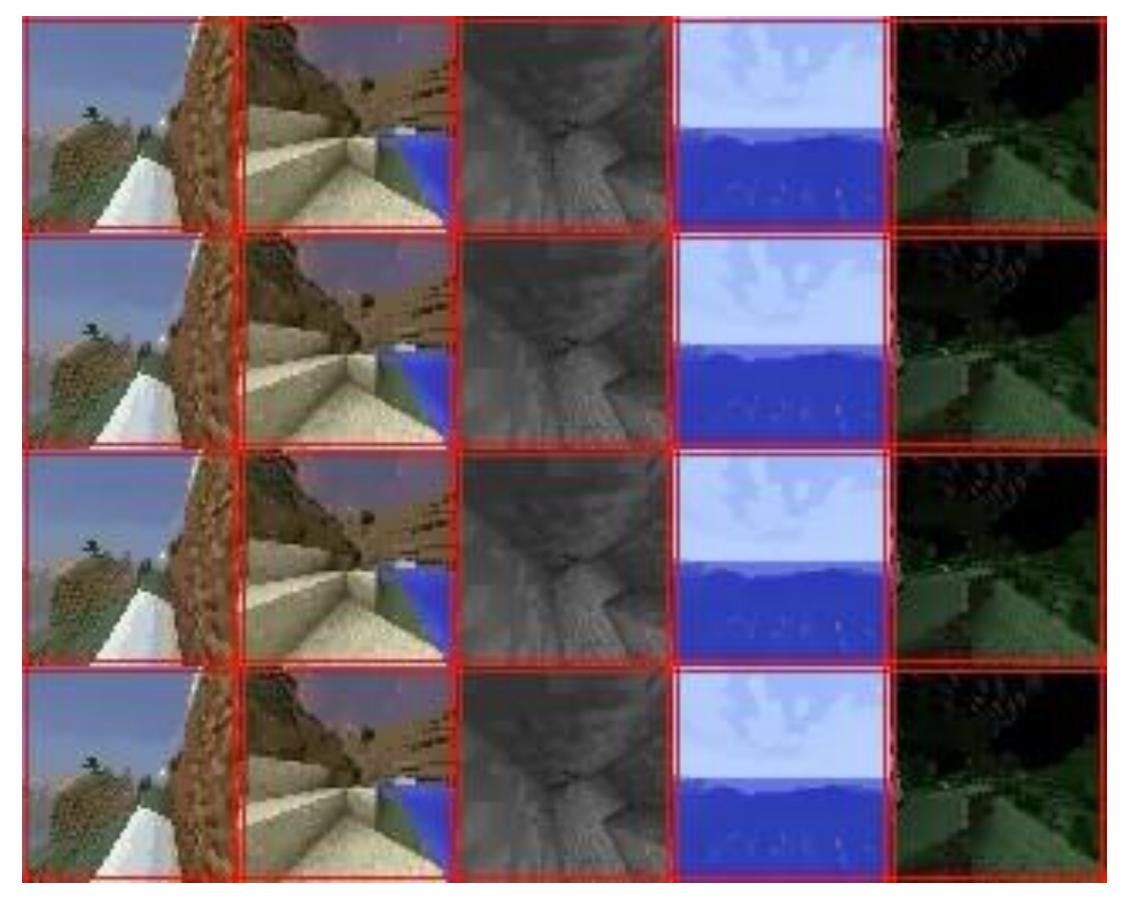
Saharia, C., Chan, W., Chang, H., Lee, C., Ho, J., Salimans, T., ... & Norouzi, M. (2022, July). Palette: Image-to-image diffusion models. In ACM SIGGRAPH 2022 Conference Proceedings (pp. 1-10).

#### VIDEO GENERATION



Ho, J., Salimans, T., Gritsenko, A., Chan, W., Norouzi, M., & Fleet, D. J. (2022). Video diffusion models. arXiv preprint arXiv:2204.03458.

#### VIDEO GENERATION



Harvey, W., Naderiparizi, S., Masrani, V., Weilbach, C., & Wood, F. (2022). Flexible Diffusion Modeling of Long Videos. arXiv preprint

## QUESTIONS?

